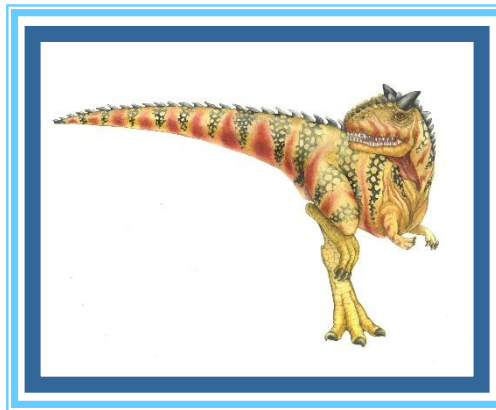


Chapter 1: Introduction





Chapter 1: Introduction

What Operating Systems Do
Computer-System Organization
Computer-System Architecture
Operating-System Operations
Resource Management
Security and Protection
Virtualization
Distributed Systems
Kernel Data Structures
Computing Environments
Free/Libre and Open-Source Operating Systems





Objectives

Describe the general organization of a computer system and the role of interrupts

Describe the components in a modern, multiprocessor computer system

Illustrate the transition from user mode to kernel mode

Discuss how operating systems are used in various computing environments

Provide examples of free and open-source operating systems





Computer System Structure

Computer system can be divided into four components:

Hardware – provides basic computing resources

- ▶ CPU, memory, I/O devices

Operating system

- ▶ Controls and coordinates use of hardware among various applications and users

Application programs – define the ways in which the system resources are used to solve the computing problems of the users

- ▶ Word processors, compilers, web browsers, database systems, video games

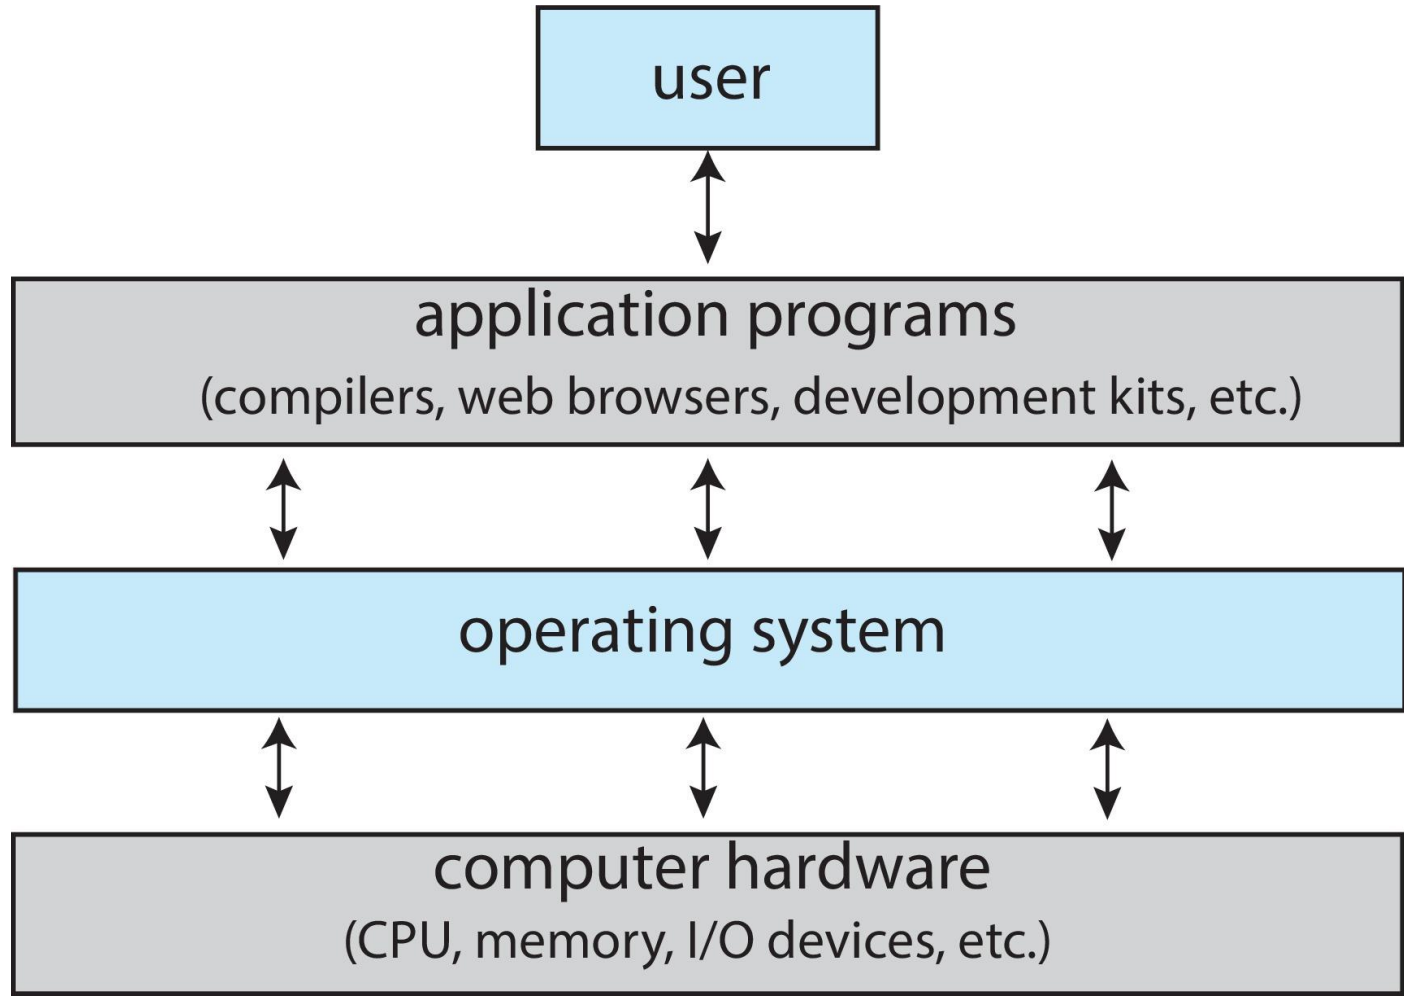
Users

- ▶ People, machines, other computers





Abstract View of Components of Computer





What Operating Systems Do

Depends on the point of view

Users want convenience, **ease of use** and **good performance**

Don't care about **resource utilization**

But shared computer such as **mainframe** or **minicomputer** must keep all users happy

Operating system is a **resource allocator** and **control program** making efficient use of HW and managing execution of user programs

Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**

Mobile devices like smartphones and tablets are resource poor, optimized for usability and battery life

Mobile user interfaces such as touch screens, voice recognition

Some computers have little or no user interface, such as embedded computers in devices and automobiles

Run primarily without user intervention





Defining Operating Systems

Term OS covers many roles

Because of myriad designs and uses of OSES

Present in toasters through ships, spacecraft, game machines, TVs and industrial control systems

Born when fixed use computers for military became more general purpose and needed resource management and program control





Operating System Definition (Cont.)

No universally accepted definition

“Everything a vendor ships when you order an operating system” is a good approximation

But varies wildly


“The one program running at all times on the computer” is the **kernel**, part of the operating system

Everything else is either

a **system program** (ships with the operating system, but not part of the kernel) , or

an **application program**, all programs not associated with the operating system

Today's OSES for general purpose and mobile computing also include **middleware** – a set of software frameworks that provide additional services to application developers such as databases, multimedia, graphics



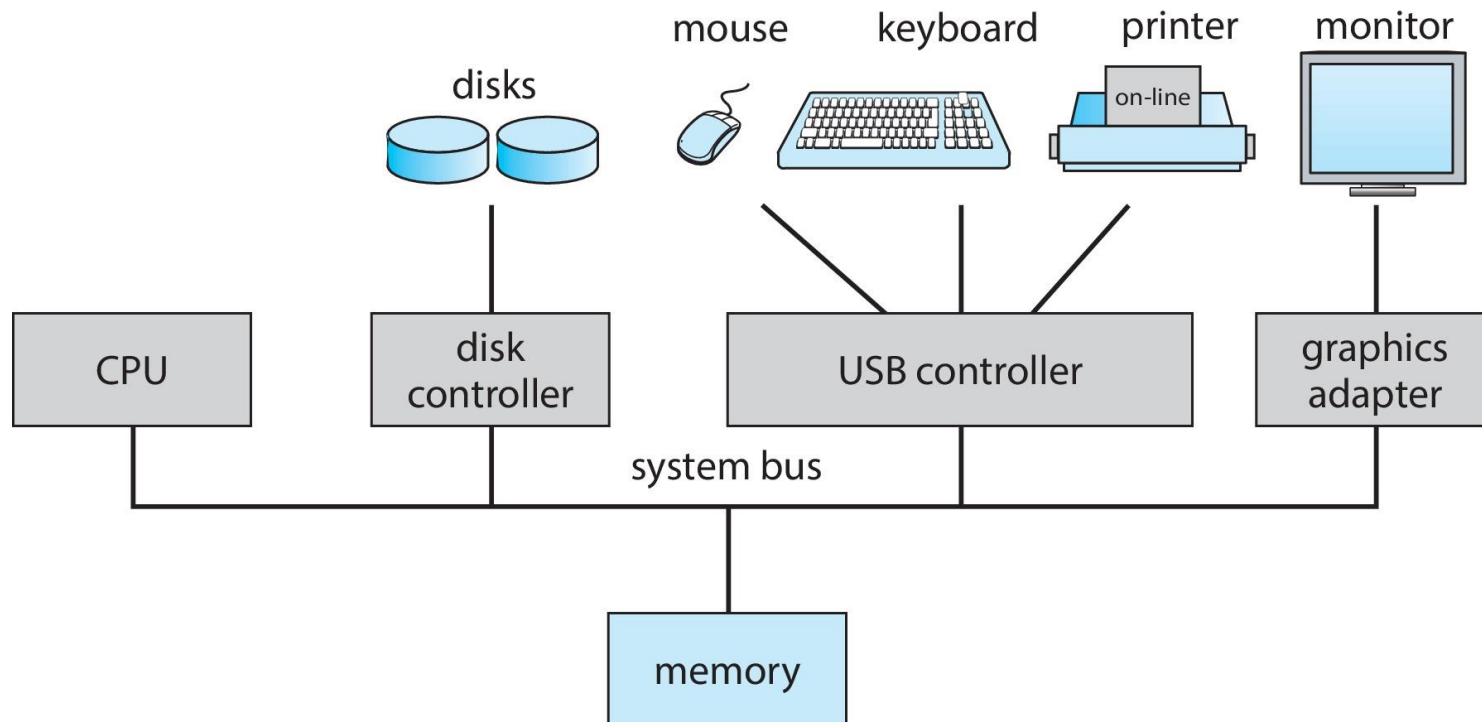


Computer System Organization

Computer-system operation

One or more CPUs, device controllers connect through common **bus** providing access to shared memory

Concurrent execution of CPUs and devices competing for memory cycles





Computer-System Operation

I/O devices and the CPU can execute concurrently

Each device controller is in charge of a particular device type

Each device controller has a local buffer

Each device controller type has an operating system **device driver** to manage it

CPU moves data from/to main memory to/from local buffers

I/O is from the device to local buffer of controller

Device controller informs CPU that it has finished its operation by causing an **interrupt**





Common Functions of Interrupts

Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines

Interrupt architecture must save the address of the interrupted instruction

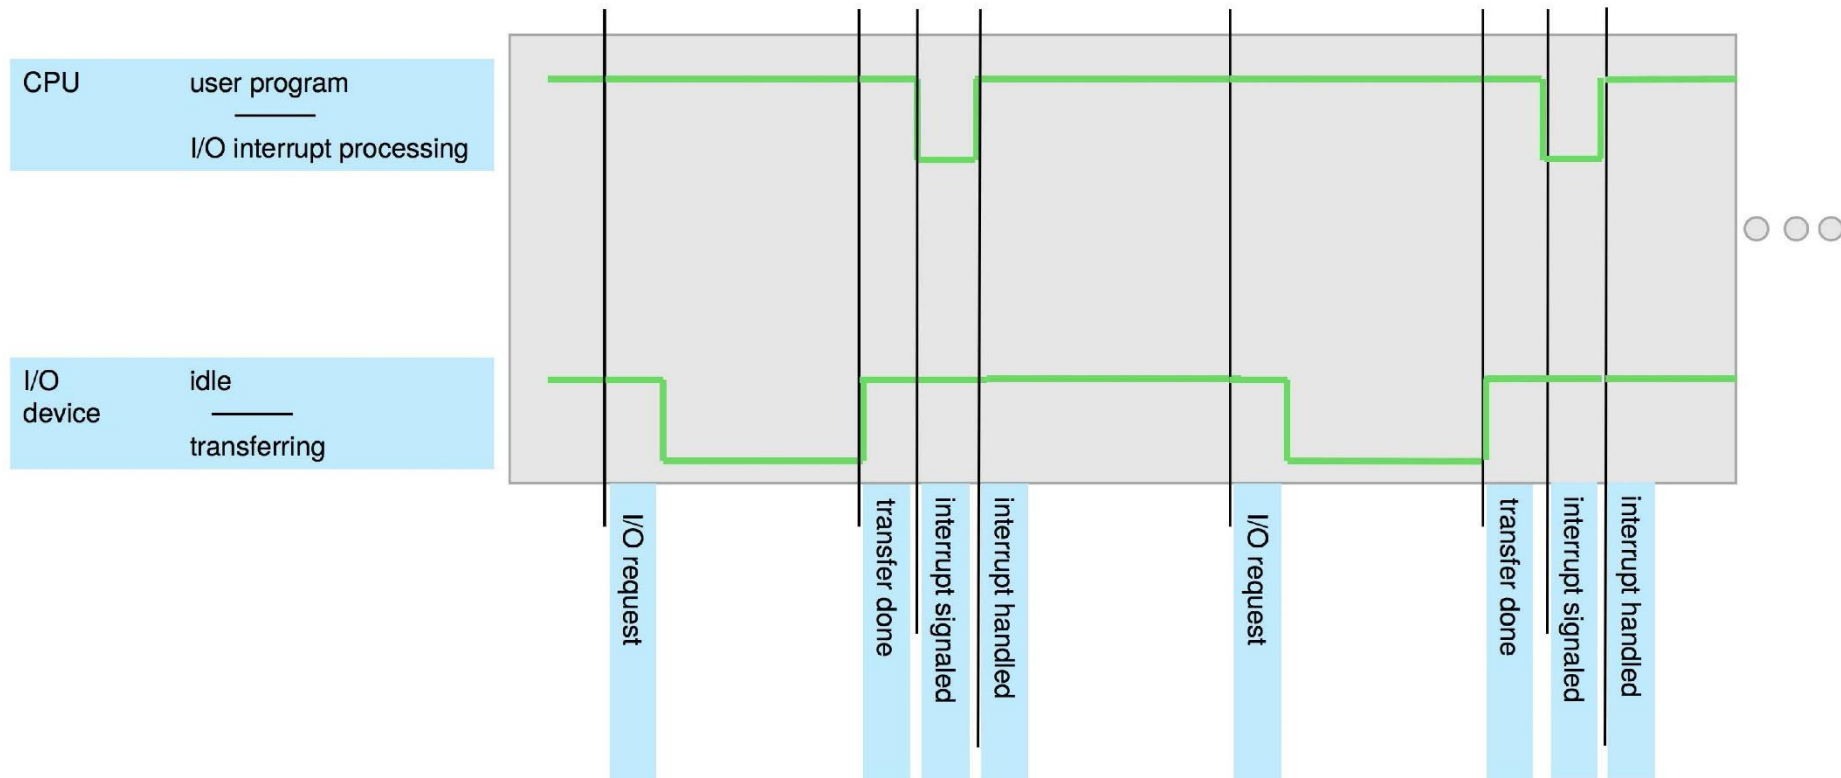
A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request

An operating system is **interrupt driven**





Interrupt Timeline





Computer Startup

bootstrap program is loaded at power-up or reboot

Typically stored in ROM or EPROM, generally known as **firmware**

Initializes all aspects of system

Loads operating system kernel and starts execution





Interrupt Handling

The operating system preserves the state of the CPU by storing registers and the program counter

Determines which type of interrupt has occurred:

polling

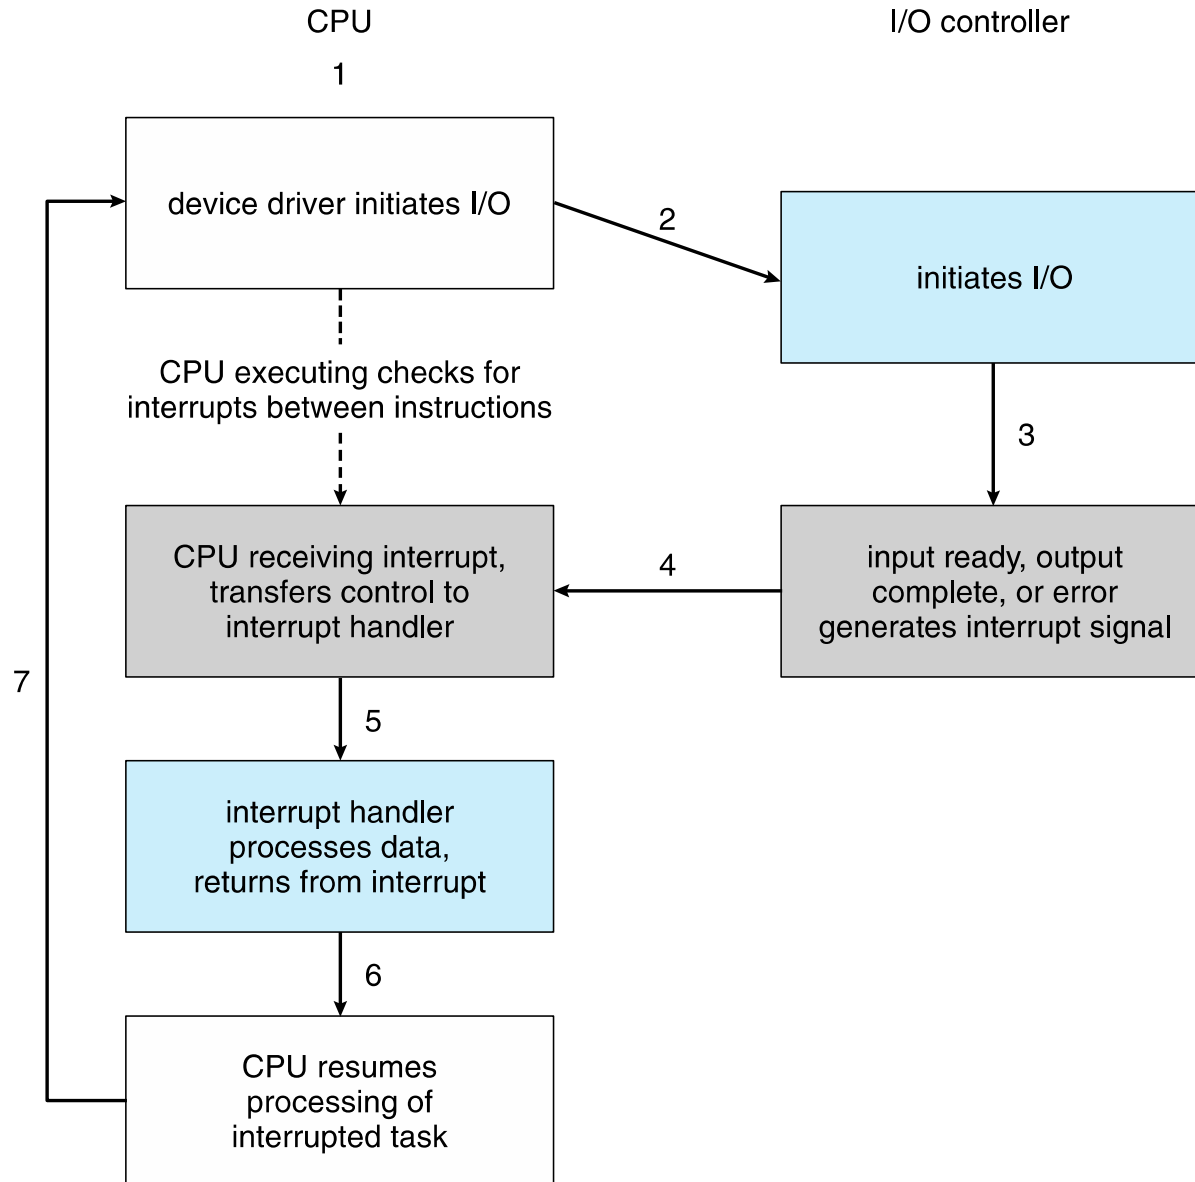
vectored interrupt system

Separate segments of code determine what action should be taken for each type of interrupt





Interrupt-drive I/O Cycle





I/O Structure

After I/O starts, control returns to user program only upon I/O completion

Wait instruction idles the CPU until the next interrupt

Wait loop (contention for memory access)

At most one I/O request is outstanding at a time, no simultaneous I/O processing

After I/O starts, control returns to user program without waiting for I/O completion

System call – request to the OS to allow user to wait for I/O completion

Device-status table contains entry for each I/O device indicating its type, address, and state

OS indexes into I/O device table to determine device status and to modify table entry to include interrupt





Storage Structure

Main memory – only large storage media that the CPU can access directly

Random access

Typically **volatile**

Typically **random-access memory** in the form of **Dynamic Random-access Memory (DRAM)**

Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity

Hard Disk Drives (HDD) – rigid metal or glass platters covered with magnetic recording material

Disk surface is logically divided into **tracks**, which are subdivided into **sectors**

The **disk controller** determines the logical interaction between the device and the computer

Non-volatile memory (NVM) devices– faster than hard disks, nonvolatile

Various technologies

Becoming more popular as capacity and performance increases, price drops






Storage Definitions and Notation Review

The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes. A **kilobyte**, or KB, is 1,024 bytes; a **megabyte**, or **MB**, is $1,024^2$ bytes; a **gigabyte**, or GB, is $1,024^3$ bytes; a **terabyte**, or **TB**, is $1,024^4$ bytes; and a **petabyte**, or **PB**, is $1,024^5$ bytes. Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).





Storage Hierarchy

Storage systems organized in hierarchy

Speed

Cost

Volatility

Caching – copying information into faster storage system; main memory can be viewed as a cache for secondary storage

Device Driver for each device controller to manage I/O

Provides uniform interface between controller and kernel

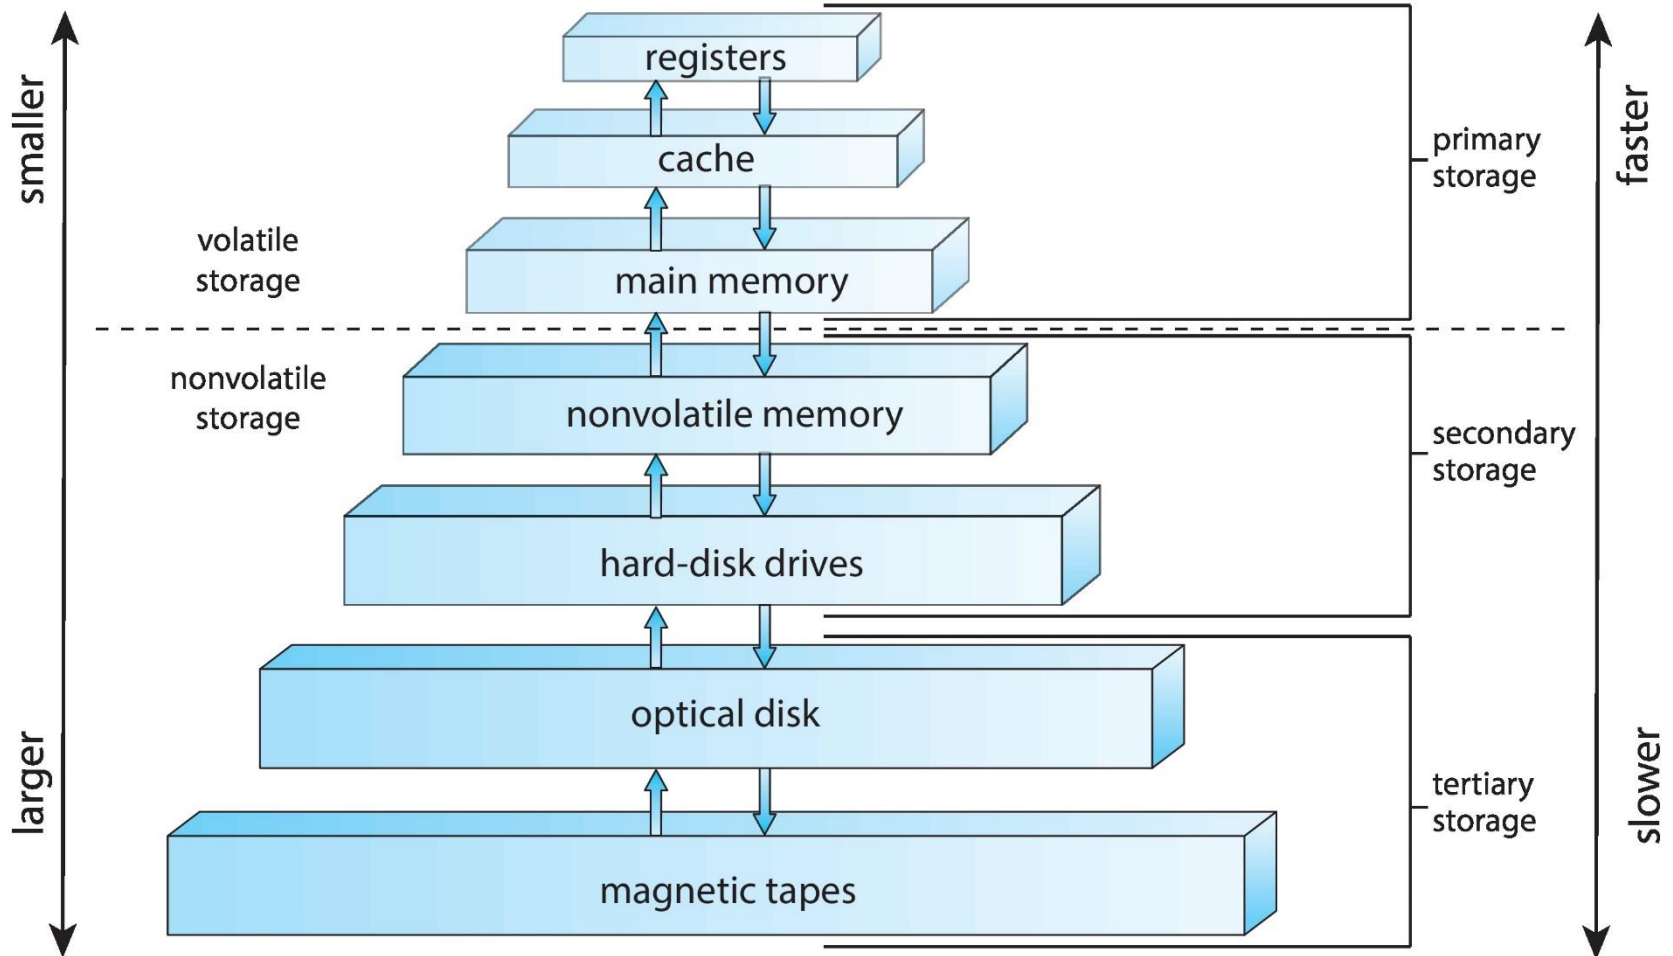




Storage-Device Hierarchy

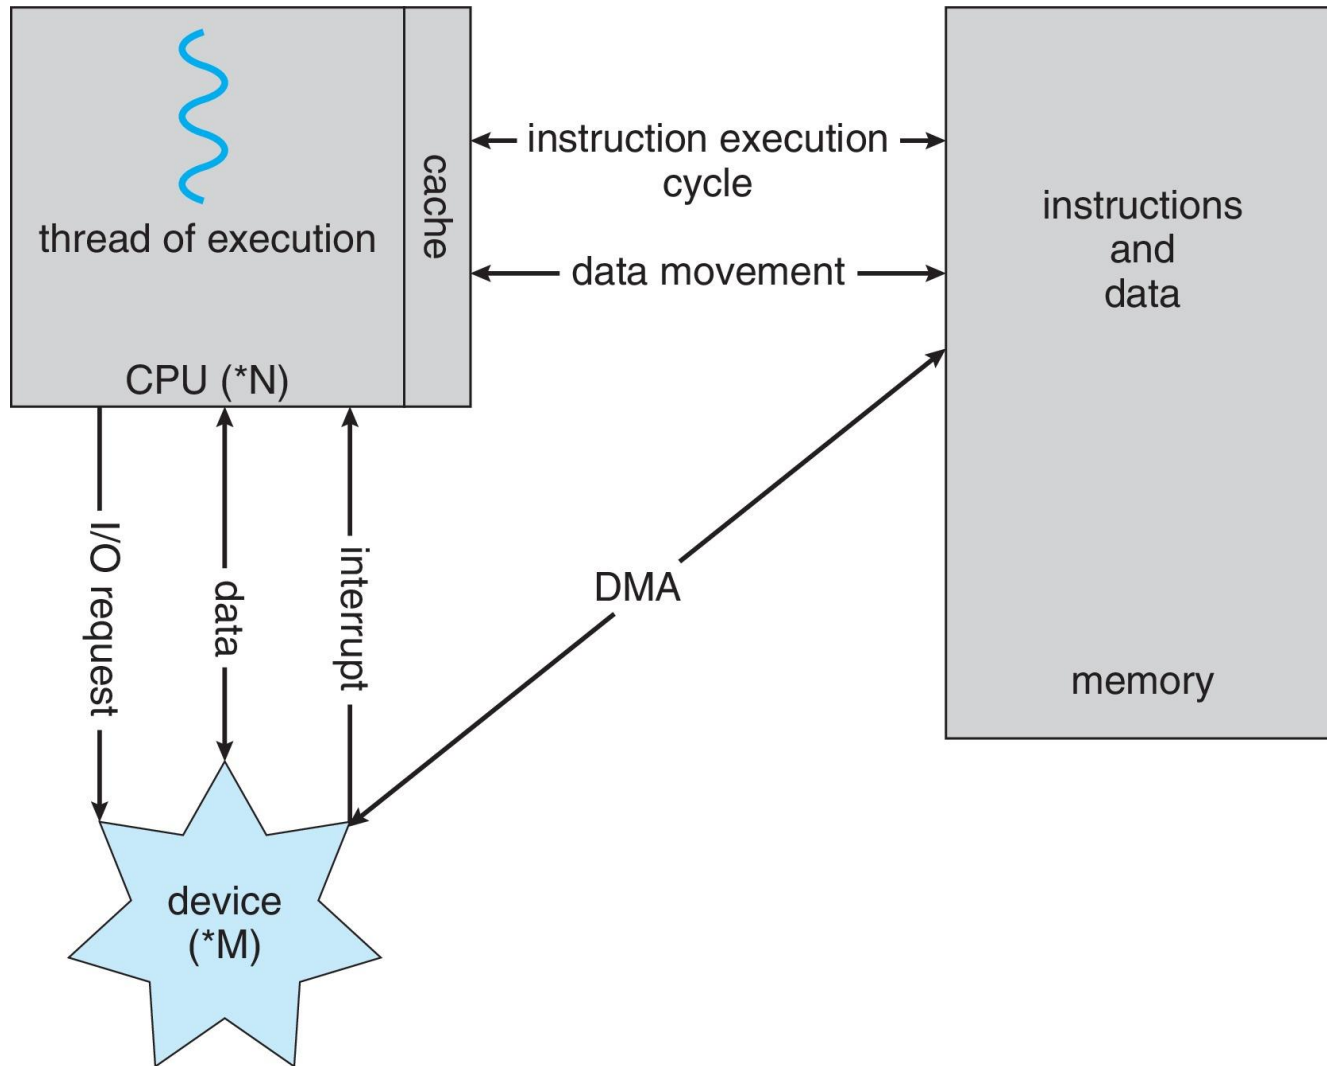
storage capacity

access time





How a Modern Computer Works



A von Neumann architecture





Direct Memory Access Structure

Used for high-speed I/O devices able to transmit information at close to memory speeds

Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention

Only one interrupt is generated per block, rather than the one interrupt per byte





Computer-System Architecture

Most systems use a single general-purpose processor

Most systems have special-purpose processors as well

Multiprocessors systems growing in use and importance

Also known as **parallel systems**, **tightly-coupled systems**

Advantages include:

1. **Increased throughput**
2. **Economy of scale**
3. **Increased reliability** – graceful degradation or fault tolerance

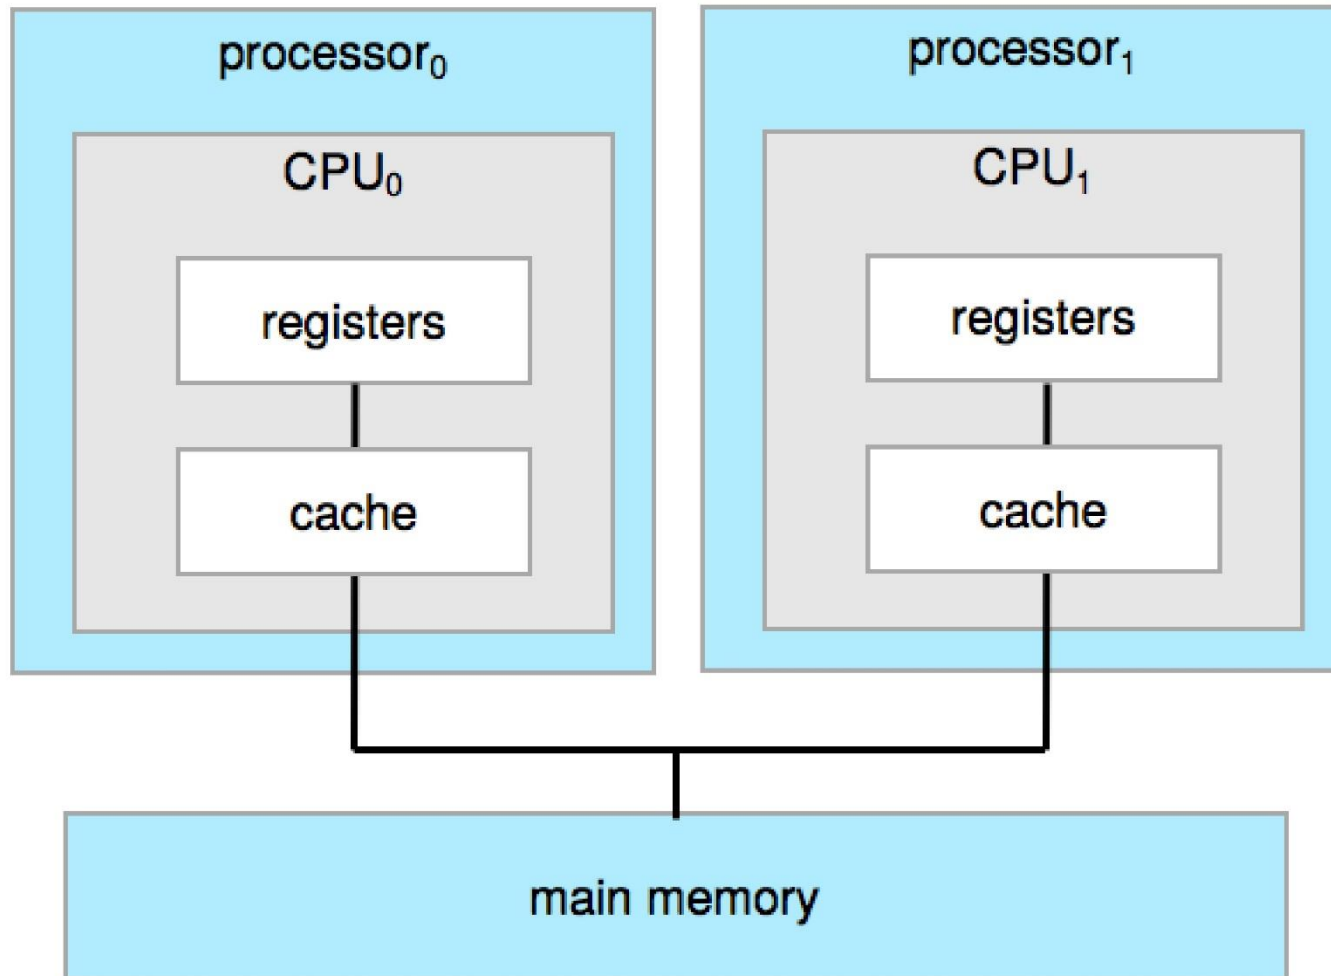
Two types:

1. **Asymmetric Multiprocessing** – each processor is assigned a specific task.
2. **Symmetric Multiprocessing** – each processor performs all tasks





Symmetric Multiprocessing Architecture







A Dual-Core Design

Multi-chip and **multicore**


Systems containing all chips

Chassis containing multiple separate systems





Non-Uniform Memory Access System





Clustered Systems

Like multiprocessor systems, but multiple systems working together

Usually sharing storage via a **storage-area network (SAN)**

Provides a **high-availability** service which survives failures

- ▶ **Asymmetric clustering** has one machine in hot-standby mode
- ▶ **Symmetric clustering** has multiple nodes running applications, monitoring each other

Some clusters are for **high-performance computing (HPC)**


- ▶ Applications must be written to use **parallelization**

Some have **distributed lock manager (DLM)** to avoid conflicting operations






Clustered Systems





PC Motherboard

Consider the desktop PC motherboard with a processor socket shown below:



This board is a fully-functioning computer, once its slots are populated. It consists of a processor socket containing a CPU, DRAM sockets, PCIe bus slots, and I/O connectors of various types. Even the lowest-cost general-purpose CPU contains multiple cores. Some motherboards contain multiple processor sockets. More advanced computers allow more than one system board, creating NUMA systems.





Operating-System Operations

Bootstrap program – simple code to initialize the system, load the kernel

Kernel loads

Starts **system daemons** (services provided outside of the kernel)

Kernel **interrupt driven** (hardware and software)

Hardware interrupt by one of the devices

Software interrupt (**exception** or **trap**):

- ▶ Software error (e.g., division by zero)
- ▶ Request for operating system service – **system call**
- ▶ Other process problems include infinite loop, processes modifying each other or the operating system





Multiprogramming and Multitasking

Multiprogramming (**Batch system**) needed for efficiency

Single user cannot keep CPU and I/O devices busy at all times

Multiprogramming organizes jobs (code and data) so CPU always has one to execute

A subset of total jobs in system is kept in memory

One job selected and run via **job scheduling**

When it has to wait (for I/O for example), OS switches to another job

Timesharing (**multitasking**) is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing

Response time should be < 1 second

Each user has at least one program executing in memory \Rightarrow **process**

If several jobs ready to run at the same time \Rightarrow **CPU scheduling**


If processes don't fit in memory, **swapping** moves them in and out to run

Virtual memory allows execution of processes not completely in memory





Memory Layout for Multiprogrammed System





Dual-mode and Multimode Operation

Dual-mode operation allows OS to protect itself and other system components

User mode and **kernel mode**

Mode bit provided by hardware

- ▶ Provides ability to distinguish when system is running user code or kernel code
- ▶ Some instructions designated as **privileged**, only executable in kernel mode
- ▶ System call changes mode to kernel, return from call resets it to user

Increasingly CPUs support multi-mode operations

i.e. **virtual machine manager (VMM)** mode for guest **VMs**





Transition from User to Kernel Mode

Timer to prevent infinite loop / process hogging resources

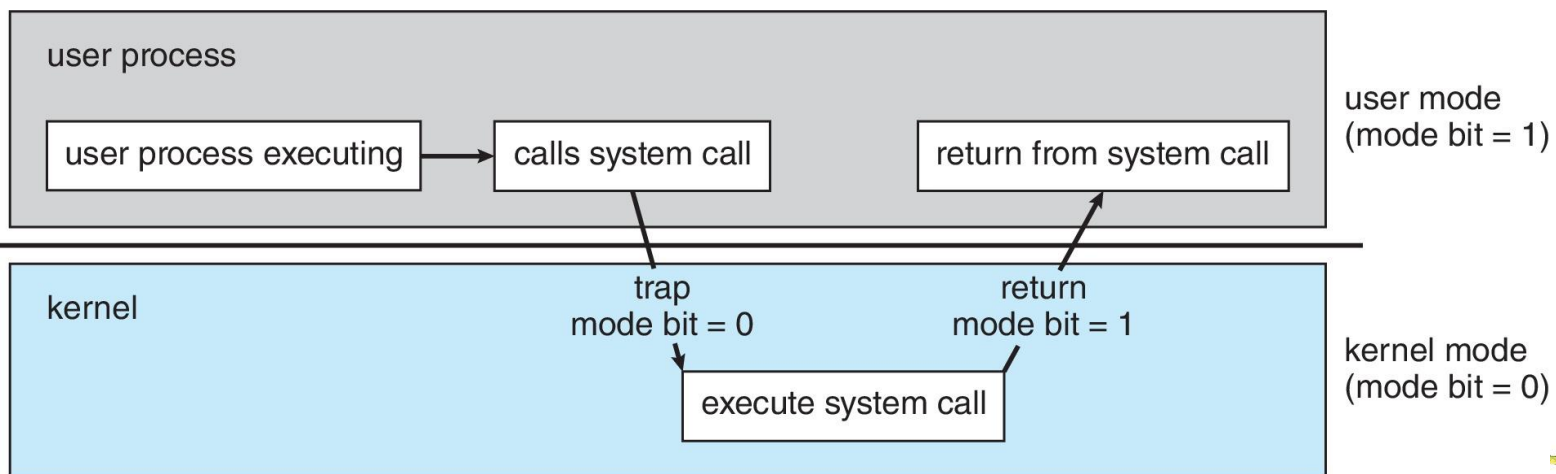
Timer is set to interrupt the computer after some time period

Keep a counter that is decremented by the physical clock

Operating system set the counter (privileged instruction)

When counter zero generate an interrupt

Set up before scheduling process to regain control or terminate program that exceeds allotted time





Process Management

A process is a program in execution. It is a unit of work within the system. Program is a ***passive entity***, process is an ***active entity***.

Process needs resources to accomplish its task

- CPU, memory, I/O, files

- Initialization data

Process termination requires reclaim of any reusable resources


Single-threaded process has one **program counter** specifying location of next instruction to execute

- Process executes instructions sequentially, one at a time, until completion

Multi-threaded process has one program counter per thread

Typically system has many processes, some user, some operating system running concurrently on one or more CPUs

- Concurrency by multiplexing the CPUs among the processes / threads





Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- Creating and deleting both user and system processes

- Suspending and resuming processes

- Providing mechanisms for process synchronization

- Providing mechanisms for process communication

- Providing mechanisms for deadlock handling





Memory Management

To execute a program all (or part) of the instructions must be in memory

All (or part) of the data that is needed by the program must be in memory

Memory management determines what is in memory and when

Optimizing CPU utilization and computer response to users

Memory management activities

Keeping track of which parts of memory are currently being used and by whom

Deciding which processes (or parts thereof) and data to move into and out of memory

Allocating and deallocating memory space as needed





File-system Management

OS provides uniform, logical view of information storage

Abstracts physical properties to logical storage unit - **file**

Each medium is controlled by device (i.e., disk drive, tape drive)

- ▶ Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)

File-System management

Files usually organized into directories

Access control on most systems to determine who can access what

OS activities include

- ▶ Creating and deleting files and directories
- ▶ Primitives to manipulate files and directories
- ▶ Mapping files onto secondary storage
- ▶ Backup files onto stable (non-volatile) storage media





Mass-Storage Management

Usually disks used to store data that does not fit in main memory or data that must be kept for a “long” period of time

Proper management is of central importance

Entire speed of computer operation hinges on disk subsystem and its algorithms

OS activities

- Mounting and unmounting

- Free-space management

- Storage allocation

- Disk scheduling

- Partitioning

- Protection

Some storage need not be fast

- Tertiary storage includes optical storage, magnetic tape

- Still must be managed – by OS or applications





Caching

Important principle, performed at many levels in a computer (in hardware, operating system, software)

Information in use copied from slower to faster storage temporarily

Faster storage (cache) checked first to determine if information is there

- If it is, information used directly from the cache (fast)

- If not, data copied to cache and used there

Cache smaller than storage being cached

- Cache management important design problem

- Cache size and replacement policy





Characteristics of Various Types of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid-state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25-0.5	0.5-25	80-250	25,000-50,000	5,000,000
Bandwidth (MB/sec)	20,000-100,000	5,000-10,000	1,000-5,000	500	20-150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape


Movement between levels of storage hierarchy can be explicit or implicit





Migration of data “A” from Disk to Register

Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache

Distributed environment situation even more complex

Several copies of a datum can exist

Various solutions covered in Chapter 19





I/O Subsystem

One purpose of OS is to hide peculiarities of hardware devices from the user

I/O subsystem responsible for

Memory management of I/O including buffering (storing data temporarily while it is being transferred), caching (storing parts of data in faster storage for performance), spooling (the overlapping of output of one job with input of other jobs)

General device-driver interface

Drivers for specific hardware devices





Protection and Security

Protection – any mechanism for controlling access of processes or users to resources defined by the OS

Security – defense of the system against internal and external attacks

Huge range, including denial-of-service, worms, viruses, identity theft, theft of service

Systems generally first distinguish among users, to determine who can do what

User identities (**user IDs**, security IDs) include name and associated number, one per user

User ID then associated with all files, processes of that user to determine access control

Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file

Privilege escalation allows user to change to effective ID with more rights





Virtualization

Allows operating systems to run applications within other OSES

Vast and growing industry

Emulation used when source CPU type different from target type
(i.e. PowerPC to Intel x86)

Generally slowest method

When computer language not compiled to native code –

Interpretation

Virtualization – OS natively compiled for CPU, running **guest**
OSes also natively compiled

Consider VMware running WinXP guests, each running
applications, all on native WinXP **host** OS

VMM (virtual machine Manager) provides virtualization services





Virtualization (cont.)

Use cases involve laptops and desktops running multiple OSES for exploration or compatibility

- Apple laptop running Mac OS X host, Windows as a guest

- Developing apps for multiple OSES without having multiple systems

- QA testing applications without having multiple systems

- Executing and managing compute environments within data centers

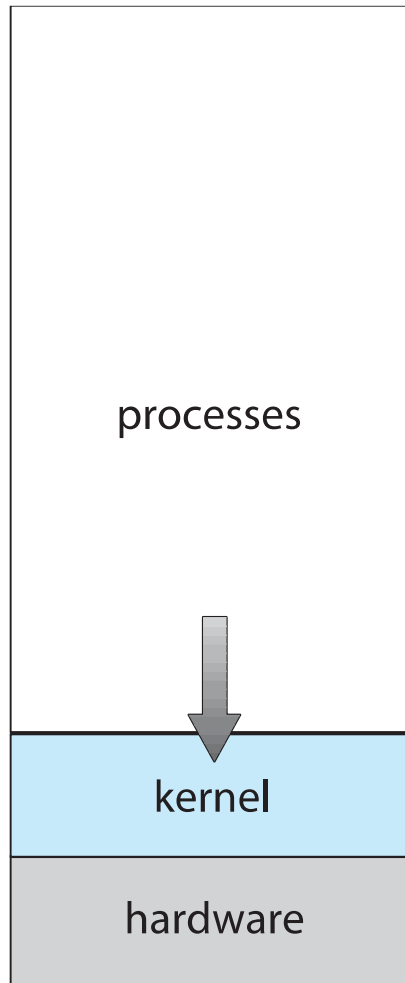
VMM can run natively, in which case they are also the host

- There is no general purpose host then (VMware ESX and Citrix XenServer)






Computing Environments - Virtualization



(a)

programming interface



(b)





Distributed Systems

Distributed computing

Collection of separate, possibly heterogeneous, systems networked together

- ▶ **Network** is a communications path, **TCP/IP** most common
 - **Local Area Network (LAN)**
 - **Wide Area Network (WAN)**
 - **Metropolitan Area Network (MAN)**
 - **Personal Area Network (PAN)**

Network Operating System provides features between systems across network

- ▶ Communication scheme allows systems to exchange messages
- ▶ Illusion of a single system







Kernel Data Structures

n Many similar to standard programming data structures



n ***Singly linked list***



n ***Doubly linked list***



n ***Circular linked list***






Kernel Data Structures

Binary search tree

left \leq right

Search performance is $O(n)$


Balanced binary search tree is $O(\lg n)$





Kernel Data Structures

Hash function can create a **hash map**



Bitmap – string of n binary digits representing the status of n items

Linux data structures defined in ***include*** files

`<linux/list.h>`, `<linux/kfifo.h>`,

`<linux/rbtree.h>`





Computing Environments - Traditional

Stand-alone general purpose machines

But blurred as most systems interconnect with others (i.e., the Internet)

Portals provide web access to internal systems

Network computers (**thin clients**) are like Web terminals

Mobile computers interconnect via **wireless networks**

Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks





Computing Environments - Mobile

Handheld smartphones, tablets, etc

What is the functional difference between them and a “traditional” laptop?

Extra feature – more OS features (GPS, gyroscope)

Allows new types of apps like ***augmented reality***

Use IEEE 802.11 wireless, or cellular data networks for connectivity

Leaders are **Apple iOS** and **Google Android**






Computing Environments – Client-Server

Client-Server Computing

Dumb terminals supplanted by smart PCs

Many systems now **servers**, responding to requests generated by **clients**

- ▶ **Compute-server system** provides an interface to client to request services (i.e., database)
- ▶ **File-server system** provides interface for clients to store and retrieve files





Computing Environments - Peer-to-Peer

Another model of distributed system

P2P does not distinguish clients and servers


Instead all nodes are considered peers

May each act as client, server or both

Node must join P2P network

- ▶ Registers its service with central lookup service on network, or
- ▶ Broadcast request for service and respond to requests for service via ***discovery protocol***

Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype





Computing Environments – Cloud Computing

Delivers computing, storage, even apps as a service across a network

Logical extension of virtualization because it uses virtualization as the base for its functionality.

Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage

Many types

Public cloud – available via Internet to anyone willing to pay

Private cloud – run by a company for the company's own use

Hybrid cloud – includes both public and private cloud components

Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)

Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)

Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)






Computing Environments – Cloud Computing

Cloud computing environments composed of traditional OSES, plus VMMs, plus cloud management tools

Internet connectivity requires security like firewalls

Load balancers spread traffic across multiple applications





Computing Environments – Real-Time Embedded Systems

Real-time embedded systems most prevalent form of computers

Vary considerable, special purpose, limited purpose OS, **real-time OS**

Use expanding

Many other special computing environments as well

Some have OSes, some perform tasks without an OS

Real-time OS has well-defined fixed time constraints

Processing ***must*** be done within constraint

Correct operation only if constraints met





Free and Open-Source Operating Systems

Operating systems made available in source-code format rather than just binary **closed-source** and **proprietary**

Counter to the **copy protection** and **Digital Rights Management (DRM)** movement

Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**

Free software and open-source software are two different ideas championed by different groups of people

▶ <http://gnu.org/philosophy/open-source-misses-the-point.html/>

Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more

Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)

Use to run guest operating systems for exploration






The Study of Operating Systems

There has never been a more interesting time to study operating systems, and it has never been easier. The open-source movement has overtaken operating systems, causing many of them to be made available in both source and binary (executable) format. The list of operating systems available in both formats includes Linux, BSD UNIX, Solaris, and part of macOS. The availability of source code allows us to study operating systems from the inside out. Questions that we could once answer only by looking at documentation or the behavior of an operating system we can now answer by examining the code itself.

Operating systems that are no longer commercially viable have been open-sourced as well, enabling us to study how systems operated in a time of fewer CPU, memory, and storage resources. An extensive but incomplete list of open-source operating-system projects is available from https://curlie.org/Computers/Software/Operating_Systems/Open_Source/

In addition, the rise of virtualization as a mainstream (and frequently free) computer function makes it possible to run many operating systems on top of one core system. For example, VMware (<http://www.vmware.com>) provides a free “player” for Windows on which hundreds of free “virtual appliances” can run. Virtualbox (<http://www.virtualbox.com>) provides a free, open-source virtual machine manager on many operating systems. Using such tools, students can try out hundreds of operating systems without dedicated hardware.

The advent of open-source operating systems has also made it easier to make the move from student to operating-system developer. With some knowledge, some effort, and an Internet connection, a student can even create a new operating-system distribution. Just a few years ago, it was difficult or impossible to get access to source code. Now, such access is limited only by how much interest, time, and disk space a student has.



End of Chapter 1

