


Chapter 3: Processes





Chapter 3: Processes

Process Concept

Process Scheduling

Operations on Processes

Interprocess Communication

IPC in Shared-Memory Systems

IPC in Message-Passing Systems

Examples of IPC Systems

Communication in Client-Server Systems





Objectives

Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system.


Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations.

Describe and contrast interprocess communication using shared memory and message passing.

Design programs that uses pipes and POSIX shared memory to perform interprocess communication.

Describe client-server communication using sockets and remote procedure calls.

Design kernel modules that interact with the Linux operating system.





Process Concept

An operating system executes a variety of programs that run as a process.

Process – a program in execution; process execution must progress in sequential fashion

Multiple parts

The program code, also called **text section**

Current activity including **program counter**, processor registers

Stack containing temporary data

- ▶ Function parameters, return addresses, local variables

Data section containing global variables

Heap containing memory dynamically allocated during run time





Process Concept (Cont.)

Program is ***passive*** entity stored on disk (**executable file**); process is ***active***

Program becomes process when executable file loaded into memory

Execution of program started via GUI mouse clicks, command line entry of its name, etc.


One program can be several processes

Consider multiple users executing the same program



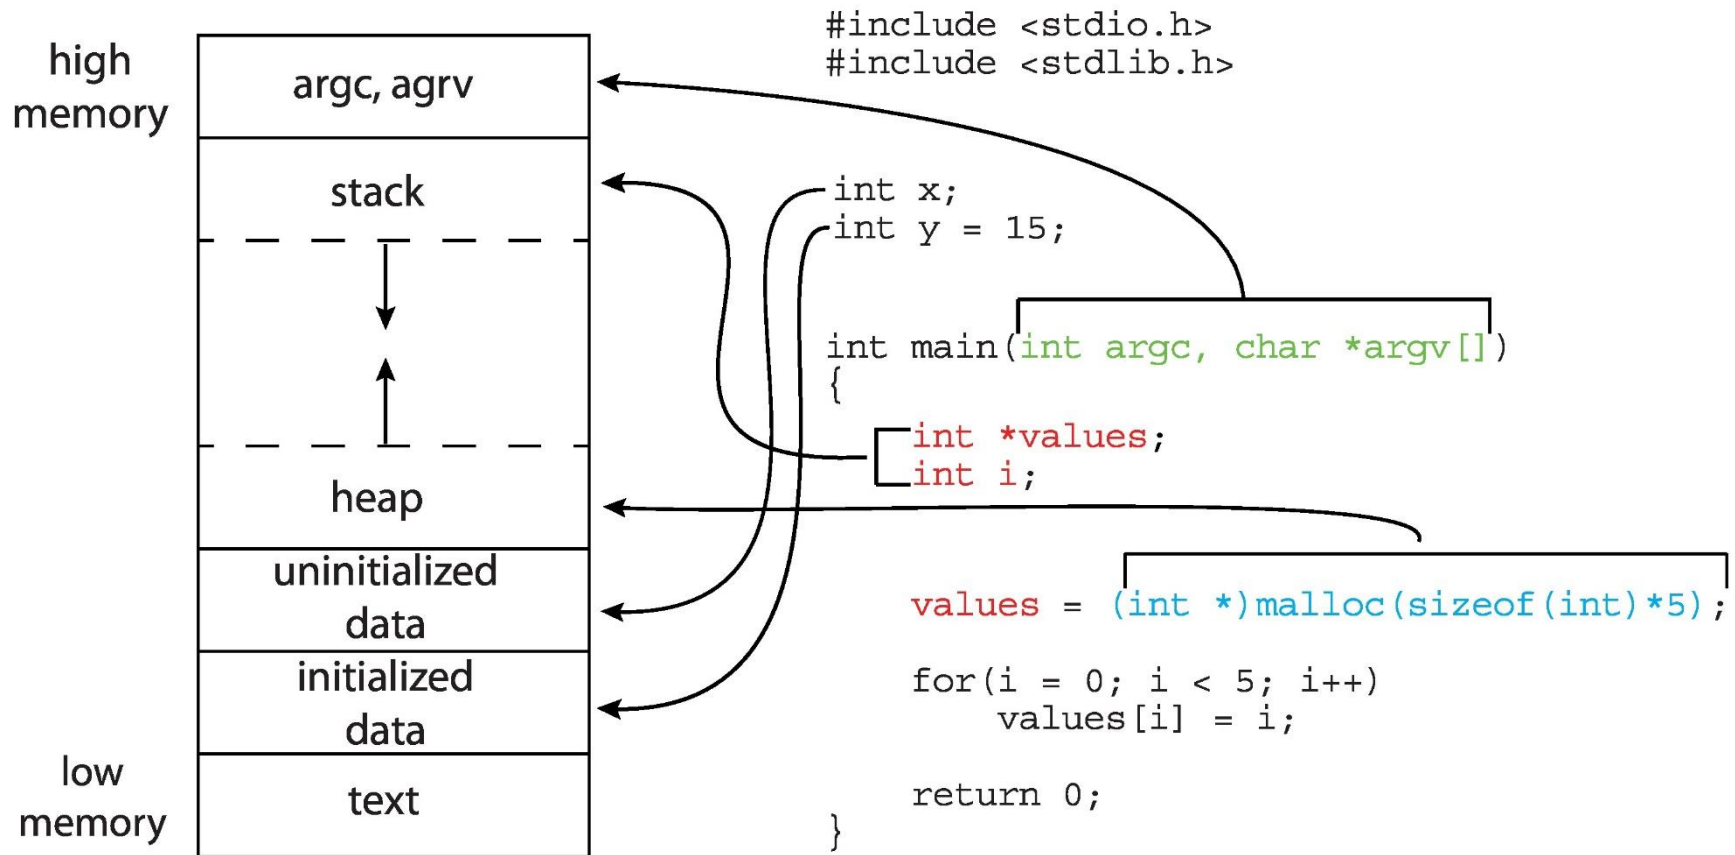


Process in Memory





Memory Layout of a C Program





Process State

As a process executes, it changes **state**

New: The process is being created

Running: Instructions are being executed

Waiting: The process is waiting for some event to occur

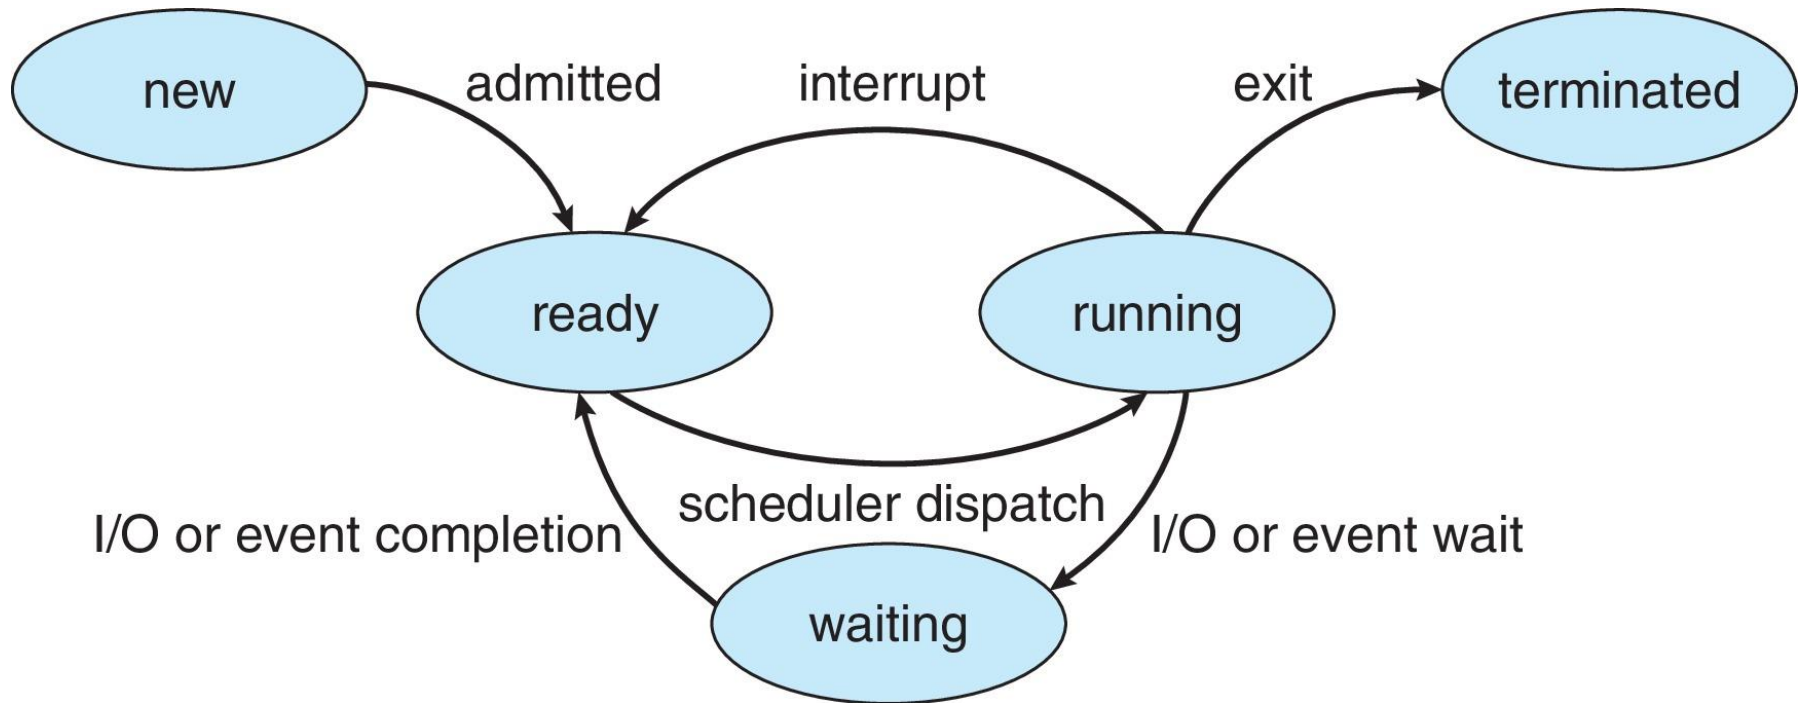
Ready: The process is waiting to be assigned to a processor

Terminated: The process has finished execution





Diagram of Process State





Process Control Block (PCB)

Information associated with each process
(also called **task control block**)

Process state – running, waiting, etc

Program counter – location of instruction to next execute

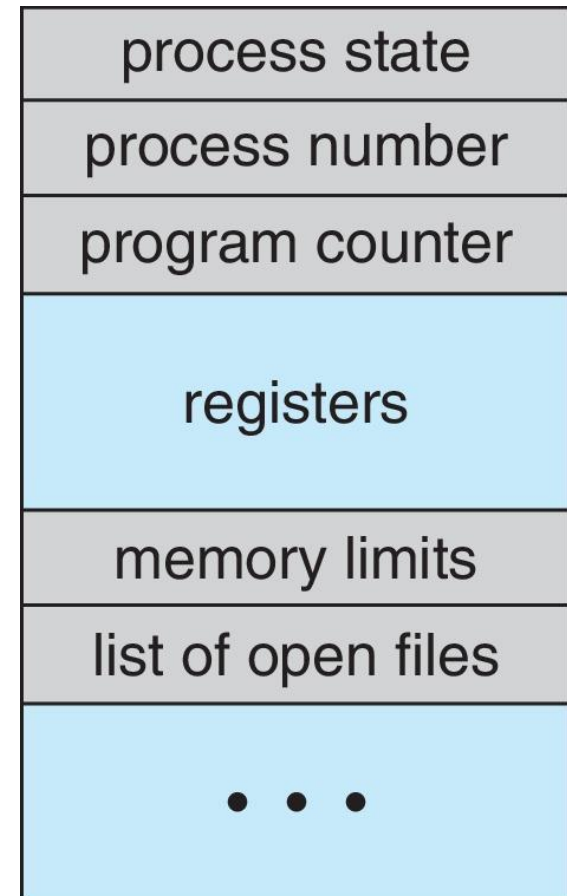
CPU registers – contents of all process-centric registers

CPU scheduling information- priorities, scheduling queue pointers

Memory-management information – memory allocated to the process

Accounting information – CPU used, clock time elapsed since start, time limits

I/O status information – I/O devices allocated to process, list of open files





Threads

So far, process has a single thread of execution

Consider having multiple program counters per process

Multiple locations can execute at once

- ▶ Multiple threads of control -> **threads**

Must then have storage for thread details, multiple program counters in PCB

Explore in detail in Chapter 4

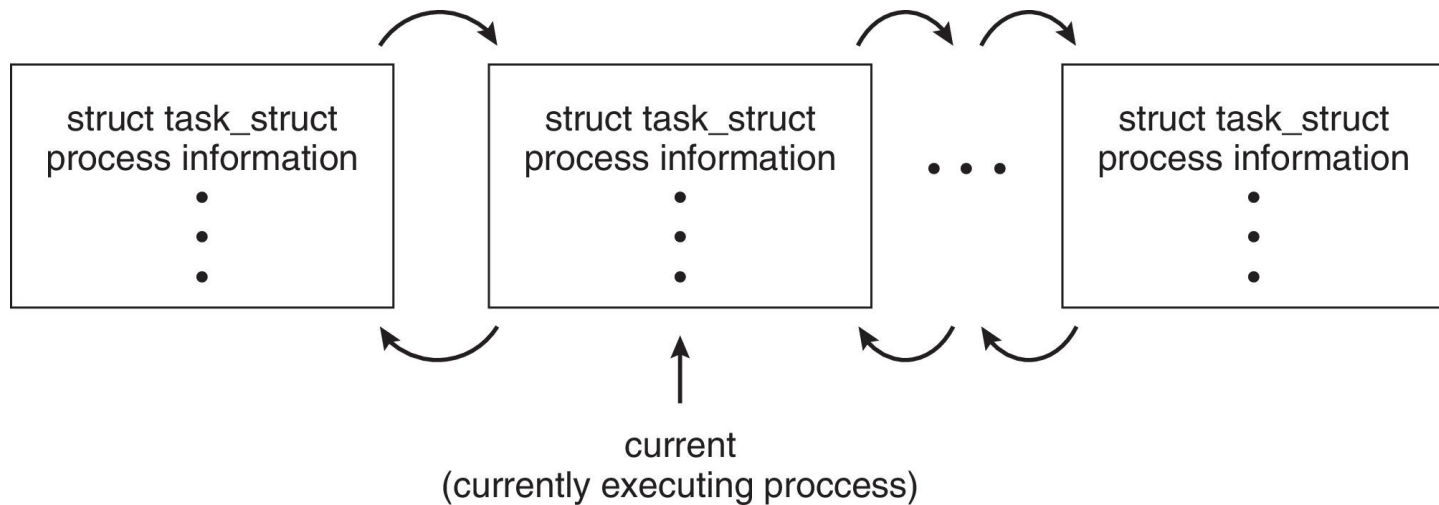




Process Representation in Linux

Represented by the C structure `task_struct`

```
pid t_pid;                /* process identifier */
long state;               /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;     /* address space of this process */
```





Process Scheduling

Maximize CPU use, quickly switch processes onto CPU core

Process scheduler selects among available processes for next execution on CPU core

Maintains **scheduling queues** of processes

Ready queue – set of all processes residing in main memory, ready and waiting to execute

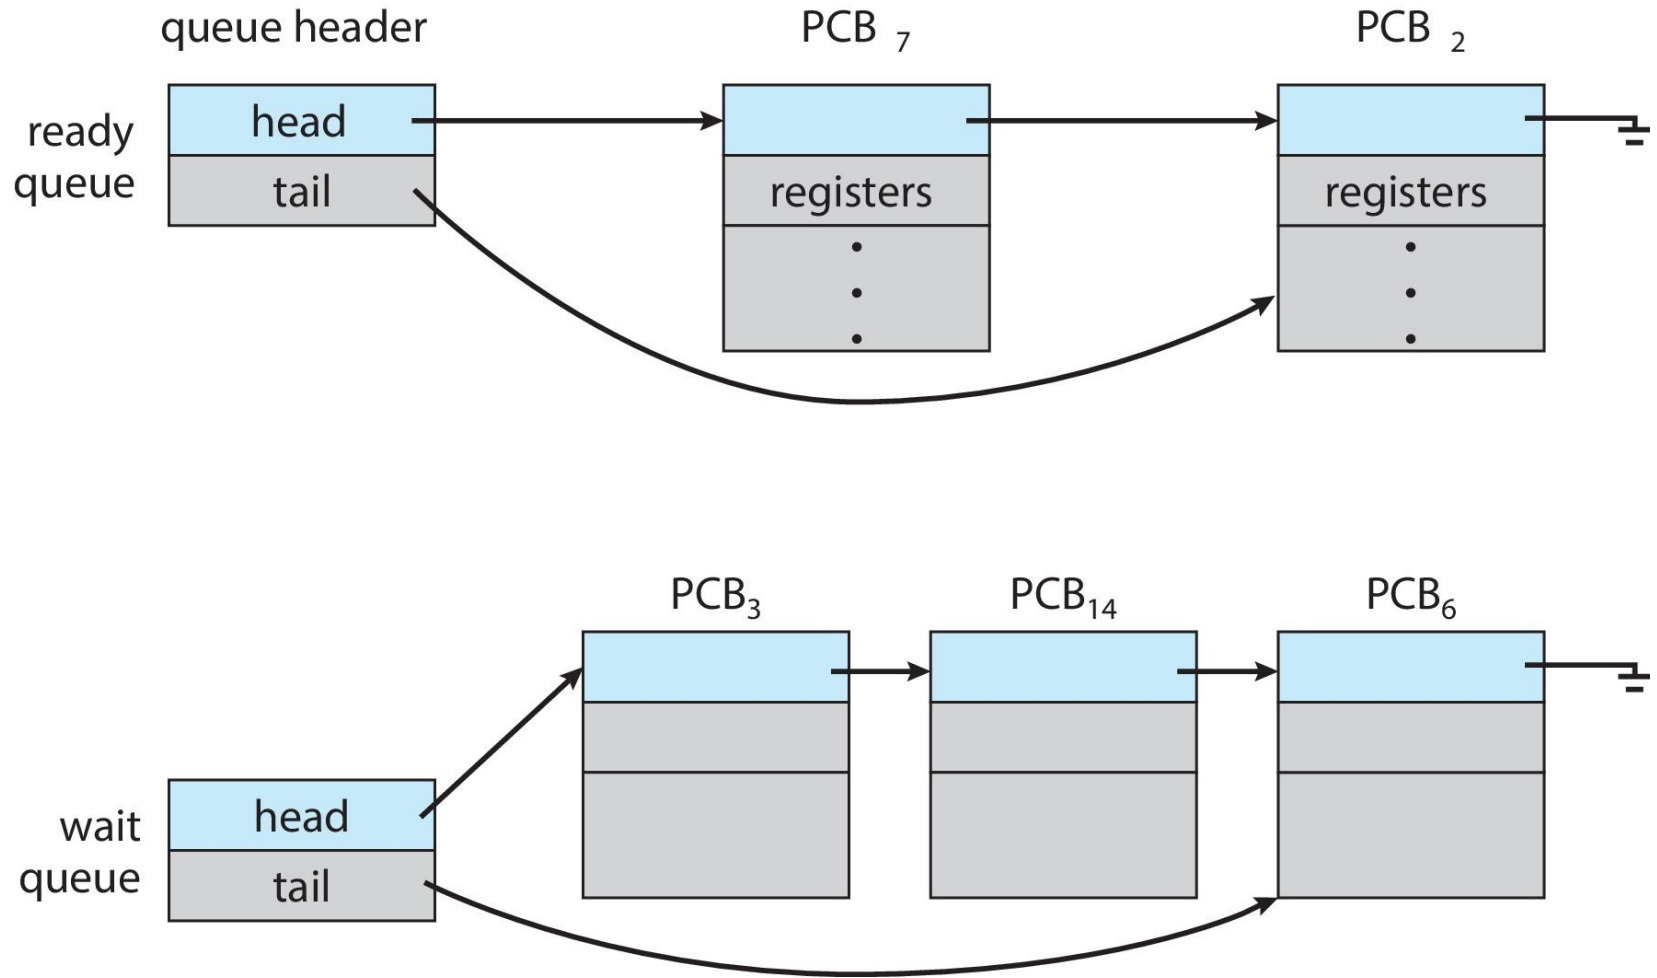
Wait queues – set of processes waiting for an event (i.e. I/O)

Processes migrate among the various queues



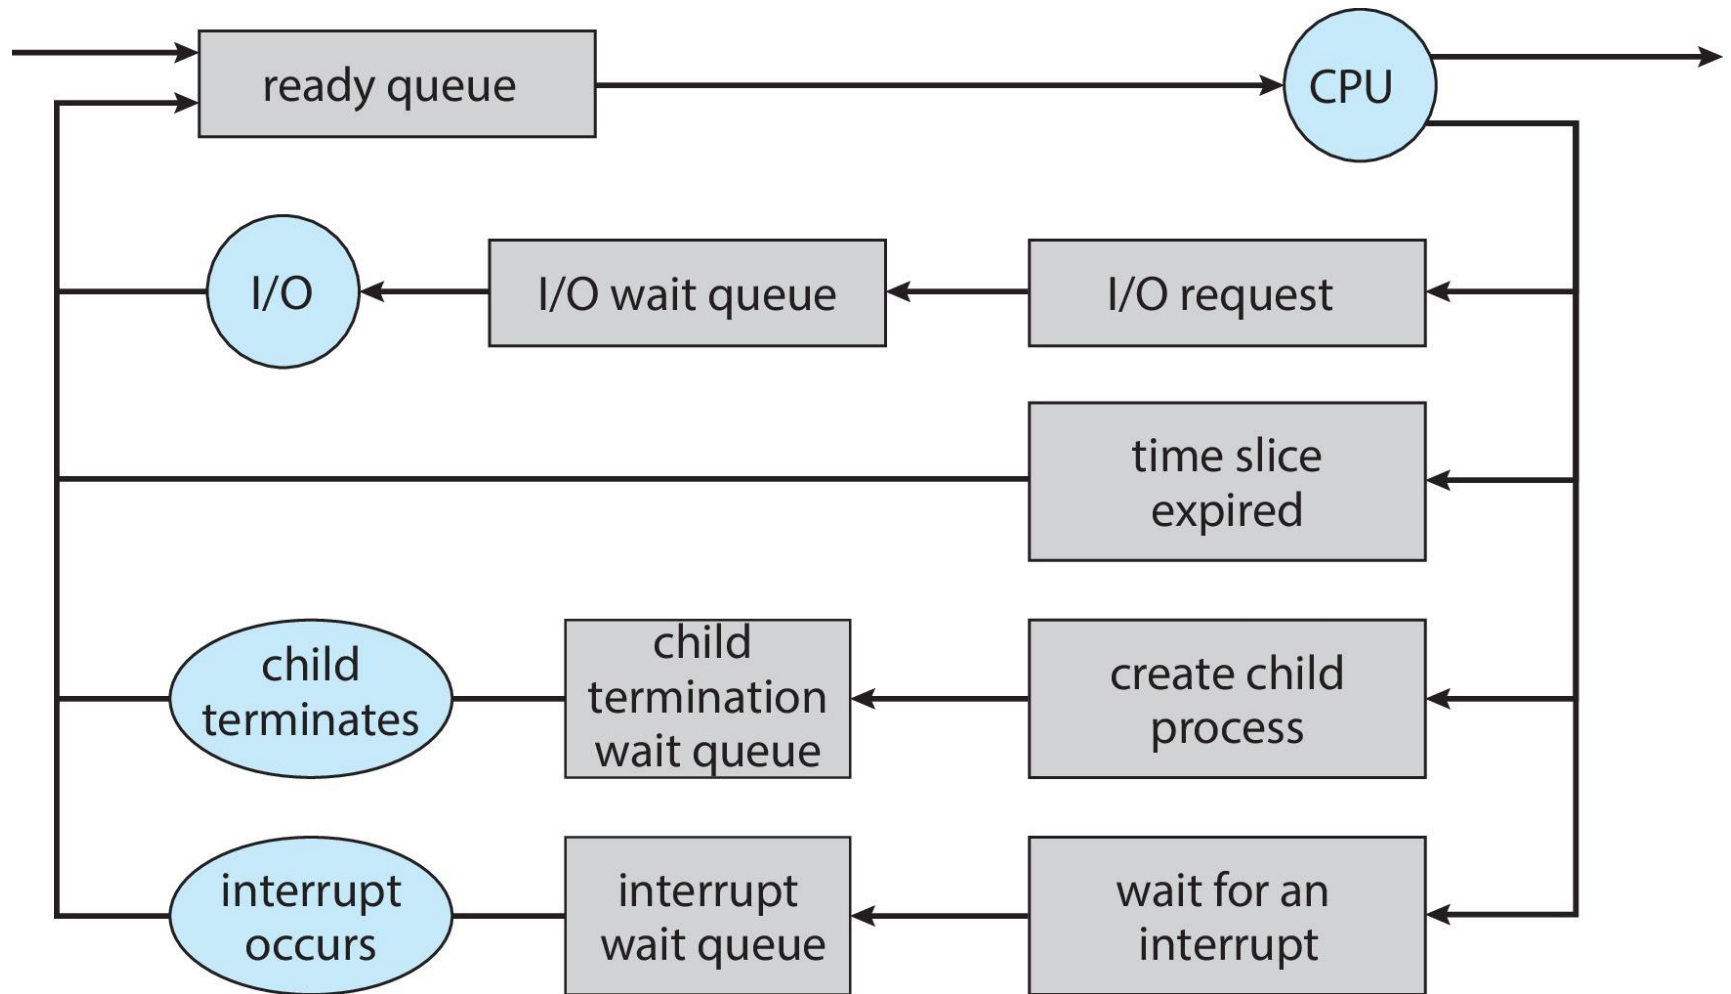


Ready and Wait Queues





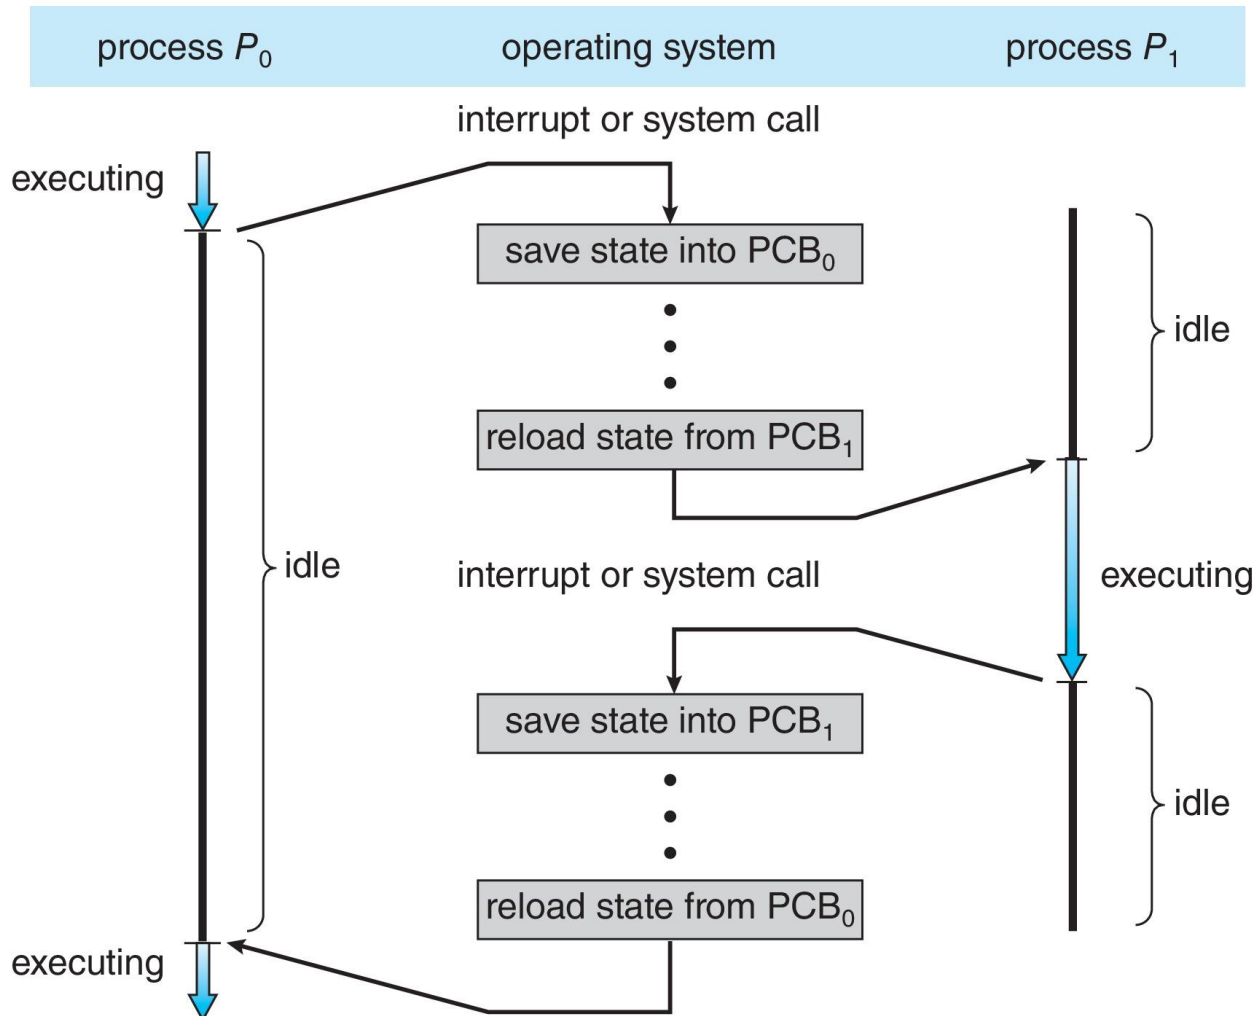
Representation of Process Scheduling





CPU Switch From Process to Process

A **context switch** occurs when the CPU switches from one process to another.





Context Switch

When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**

Context of a process represented in the PCB

Context-switch time is overhead; the system does no useful work while switching

The more complex the OS and the PCB → the longer the context switch

Time dependent on hardware support

Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once





Multitasking in Mobile Systems

Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended

Due to screen real estate, user interface limits iOS provides for a

- Single **foreground** process- controlled via user interface

- Multiple **background** processes— in memory, running, but not on the display, and with limits

- Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback

Android runs foreground and background, with fewer limits

- Background process uses a **service** to perform tasks

- Service can keep running even if background process is suspended

- Service has no user interface, small memory use





Operations on Processes

System must provide mechanisms for:

process creation

process termination





Process Creation

Parent process create **children** processes, which, in turn create other processes, forming a **tree** of processes

Generally, process identified and managed via a **process identifier (pid)**

Resource sharing options

- Parent and children share all resources

- Children share subset of parent's resources

- Parent and child share no resources

Execution options

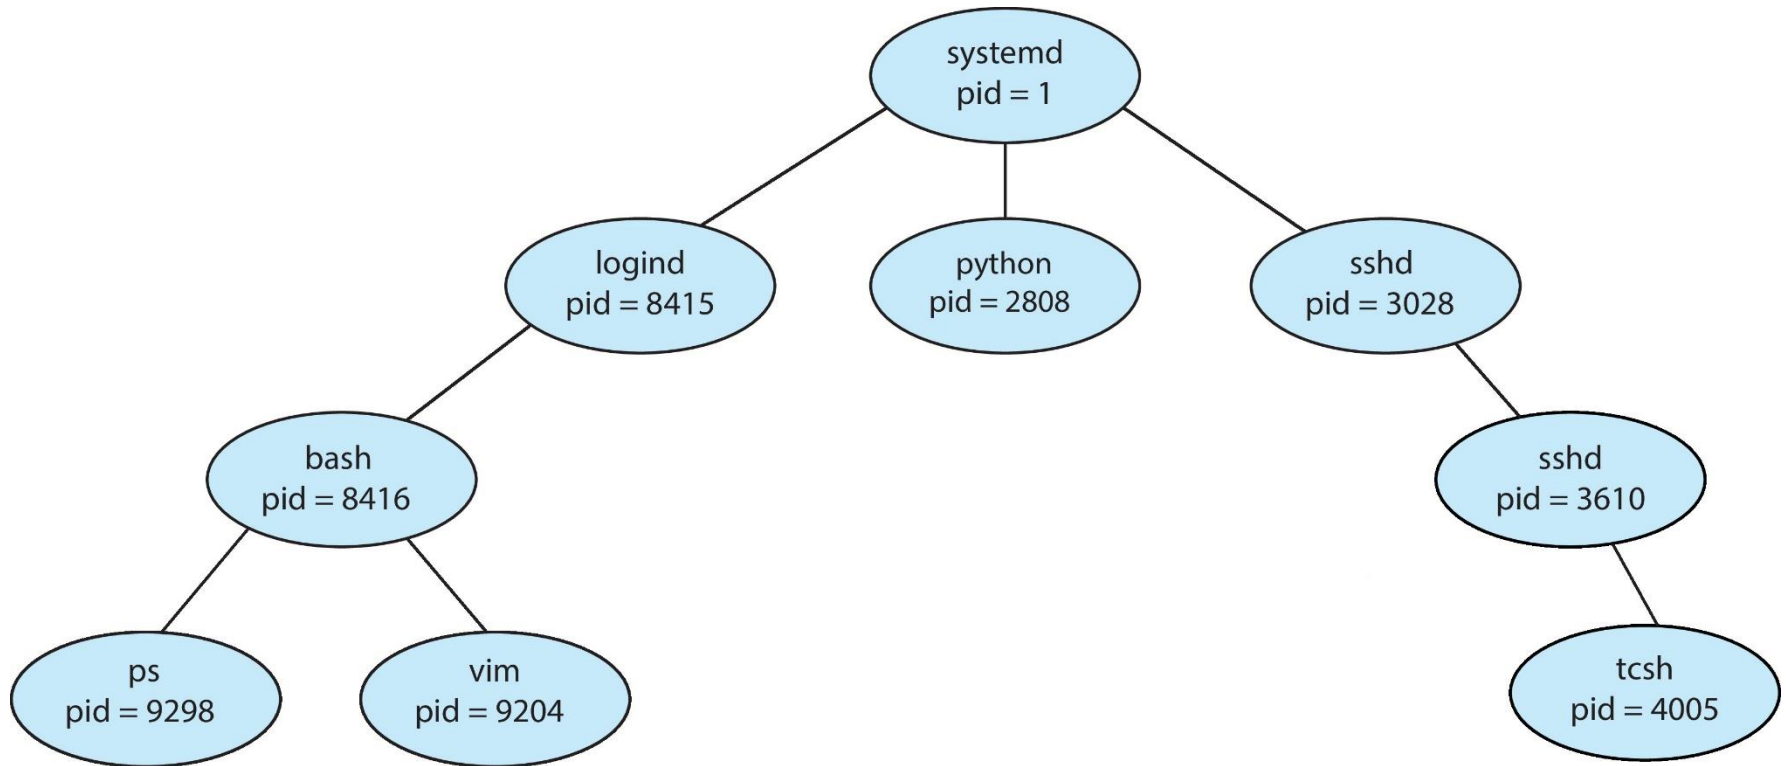
- Parent and children execute concurrently

- Parent waits until children terminate





A Tree of Processes in Linux





Process Creation (Cont.)

Address space

Child duplicate of parent

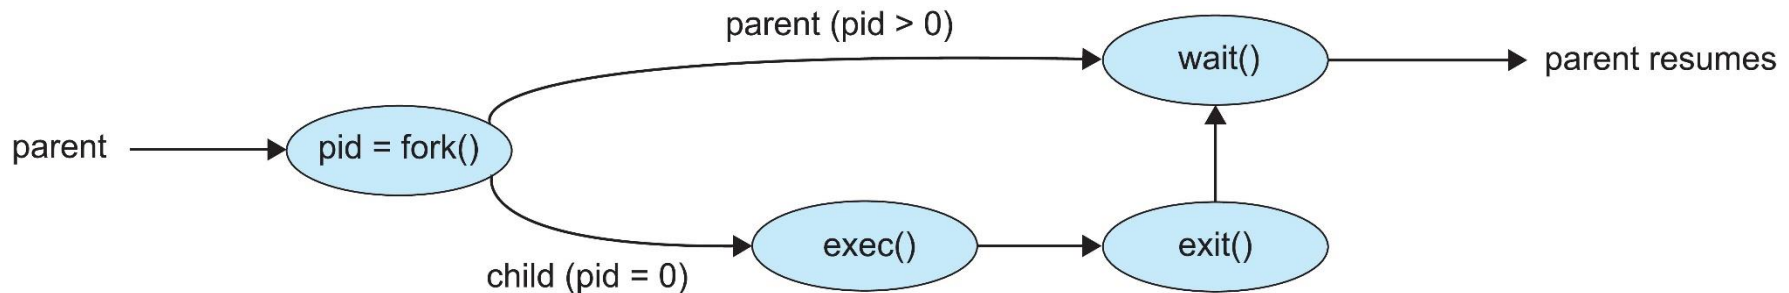
Child has a program loaded into it

UNIX examples

fork() system call creates new process

exec() system call used after a **fork()** to replace the process' memory space with a new program

Parent process calls **wait()** for the child to terminate





C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```





Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```





Process Termination

Process executes last statement and then asks the operating system to delete it using the **exit()** system call.

Returns status data from child to parent (via **wait()**)

Process' resources are deallocated by operating system

Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:

Child has exceeded allocated resources

Task assigned to child is no longer required

The parent is exiting and the operating systems does not allow a child to continue if its parent terminates





Process Termination

Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.

cascading termination. All children, grandchildren, etc. are terminated.

The termination is initiated by the operating system.

The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

If no parent waiting (did not invoke **wait()**) process is a **zombie**

If parent terminated without invoking **wait**, process is an **orphan**





Android Process Importance Hierarchy

Mobile operating systems often have to terminate processes to reclaim system resources such as memory. From **most** to **least** important:

- Foreground process
- Visible process
- Service process
- Background process
- Empty process

Android will begin terminating processes that are least important.





Multiprocess Architecture – Chrome Browser

Many web browsers ran as single process (some still do)

If one web site causes trouble, entire browser can hang or crash


Google Chrome Browser is multiprocess with 3 different types of processes:

Browser process manages user interface, disk and network I/O

Renderer process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened

- ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits

Plug-in process for each type of plug-in



Each tab represents a separate process.





Interprocess Communication

Processes within a system may be *independent* or *cooperating*

Cooperating process can affect or be affected by other processes, including sharing data

Reasons for cooperating processes:

- Information sharing

- Computation speedup

- Modularity

- Convenience

Cooperating processes need **interprocess communication (IPC)**

Two models of IPC

- Shared memory**


- Message passing**





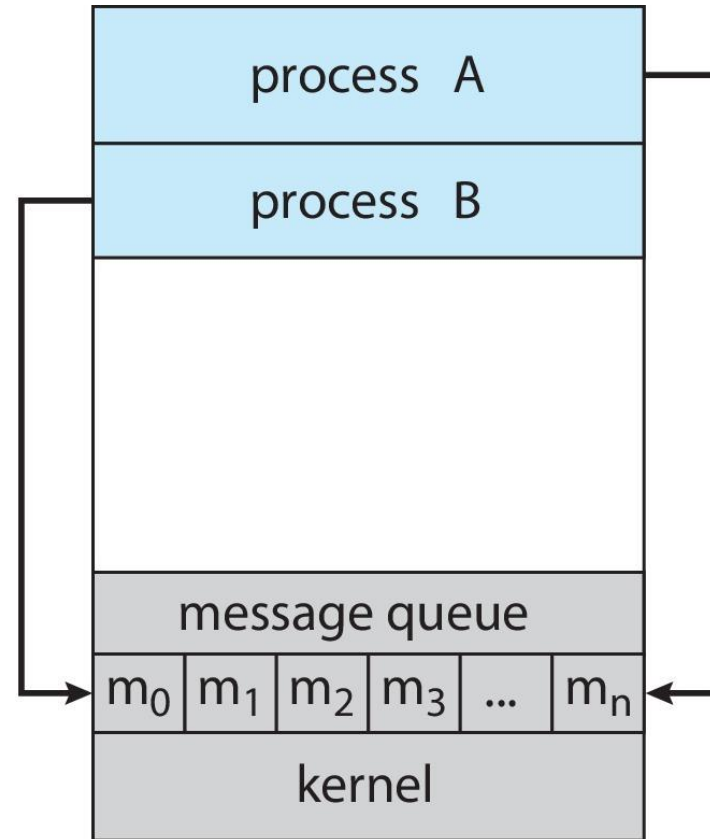
Communications Models

(a) Shared memory.



(a)

(b) Message passing.



(b)





Cooperating Processes

Independent process cannot affect or be affected by the execution of another process

Cooperating process can affect or be affected by the execution of another process

Advantages of process cooperation

- Information sharing

- Computation speed-up

- Modularity

- Convenience





Producer-Consumer Problem

Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process

unbounded-buffer places no practical limit on the size of the buffer

bounded-buffer assumes that there is a fixed buffer size





Interprocess Communication – Shared Memory

An area of memory shared among the processes that wish to communicate

The communication is under the control of the users processes not the operating system.

Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.

Synchronization is discussed in great details in Chapters 6 & 7.





Bounded-Buffer – Shared-Memory Solution

Shared data

```
#define BUFFER_SIZE 10  
typedef struct {  
    . . .  
} item;  
  
item buffer[BUFFER_SIZE];  
int in = 0;  
int out = 0;
```

Solution is correct, but can only use **BUFFER_SIZE-1** elements






Producer Process – Shared Memory

```
item next_produced;

while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```





Consumer Process – Shared Memory

```
item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```





Interprocess Communication – Message Passing

Mechanism for processes to communicate and to synchronize their actions

Message system – processes communicate with each other without resorting to shared variables

IPC facility provides two operations:

send(*message*)

receive(*message*)

The *message* size is either fixed or variable





Message Passing (Cont.)

If processes P and Q wish to communicate, they need to:

Establish a ***communication link*** between them

Exchange messages via send/receive

Implementation issues:

How are links established?

Can a link be associated with more than two processes?

How many links can there be between every pair of communicating processes?

What is the capacity of a link?

Is the size of a message that the link can accommodate fixed or variable?

Is a link unidirectional or bi-directional?





Message Passing (Cont.)

Implementation of communication link

Physical:

- ▶ Shared memory
- ▶ Hardware bus
- ▶ Network

Logical:

- ▶ Direct or indirect
- ▶ Synchronous or asynchronous
- ▶ Automatic or explicit buffering





Direct Communication

Processes must name each other explicitly:

send (P , $message$) – send a message to process P

receive(Q , $message$) – receive a message from process Q

Properties of communication link

Links are established automatically

A link is associated with exactly one pair of communicating processes

Between each pair there exists exactly one link

The link may be unidirectional, but is usually bi-directional





Indirect Communication

Messages are directed and received from mailboxes (also referred to as ports)

Each mailbox has a unique id

Processes can communicate only if they share a mailbox

Properties of communication link

Link established only if processes share a common mailbox

A link may be associated with many processes

Each pair of processes may share several communication links

Link may be unidirectional or bi-directional





Indirect Communication

Operations

- create a new mailbox (port)

- send and receive messages through mailbox

- destroy a mailbox

Primitives are defined as:

send(*A, message*) – send a message to mailbox *A*

receive(*A, message*) – receive a message from mailbox *A*





Indirect Communication

Mailbox sharing

P_1 , P_2 , and P_3 share mailbox A

P_1 sends; P_2 and P_3 receive

Who gets the message?

Solutions

Allow a link to be associated with at most two processes

Allow only one process at a time to execute a receive operation

Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.





Synchronization

Message passing may be either blocking or non-blocking

Blocking is considered **synchronous**

Blocking send -- the sender is blocked until the message is received

Blocking receive -- the receiver is blocked until a message is available

Non-blocking is considered **asynchronous**

Non-blocking send -- the sender sends the message and continue


Non-blocking receive -- the receiver receives:

A valid message, or

Null message

Different combinations possible

If both send and receive are blocking, we have a **rendezvous**





Producer – Shared Memory

```
message next_produced;  
  
while (true) {  
    /* produce an item in next_produced */  
    send(next_produced) ;  
}
```





Consumer– Shared Memory

```
message next_consumed;  
  
while (true) {  
    receive(next_consumed)  
  
    /* consume the item in next_consumed */  
}
```





Buffering

Queue of messages attached to the link.

Implemented in one of three ways

1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
2. Bounded capacity – finite length of n messages
Sender must wait if link full
3. Unbounded capacity – infinite length
Sender never waits





Examples of IPC Systems - POSIX

POSIX Shared Memory

Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR,  
0666) ;
```

Also used to open an existing segment

Set the size of the object

```
ftruncate(shm_fd, 4096) ;
```

Use `mmap()` to memory-map a file pointer to the shared memory object

Reading and writing to shared memory is done by using the pointer returned by `mmap()`.





IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```





IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```





Examples of IPC Systems - Mach

Mach communication is message based

Even system calls are messages

Each task gets two ports at creation- Kernel and Notify

Messages are sent and received using the `mach_msg()` function

Ports needed for communication, created via `mach_port_allocate()`

Send and receive are flexible, for example four options if mailbox full:

- ▶ Wait indefinitely
- ▶ Wait at most n milliseconds
- ▶ Return immediately
- ▶ Temporarily cache a message





Mach Messages

```
#include<mach/mach.h>
```

```
struct message {  
    mach_msg_header_t header;  
    int data;  
};
```

```
mach_port_t client;  
mach_port_t server;
```





Mach Message Passing - Client

```
/* Client Code */

struct message message;

// construct the header
message.header.msgh_size = sizeof(message);
message.header.msgh_remote_port = server;
message.header.msgh_local_port = client;

// send the message
mach_msg(&message.header, // message header
        MACH_SEND_MSG, // sending a message
        sizeof(message), // size of message sent
        0, // maximum size of received message - unnecessary
        MACH_PORT_NULL, // name of receive port - unnecessary
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);
```





Mach Message Passing - Server

```
/* Server Code */

struct message message;

// receive the message
mach_msg(&message.header, // message header
MACH_RCV_MSG, // sending a message
0, // size of message sent
sizeof(message), // maximum size of received message
server, // name of receive port
MACH_MSG_TIMEOUT_NONE, // no time outs
MACH_PORT_NULL // no notify port
);
```





Examples of IPC Systems – Windows

Message-passing centric via **advanced local procedure call (LPC)** facility

Only works between processes on the same system

Uses ports (like mailboxes) to establish and maintain communication channels

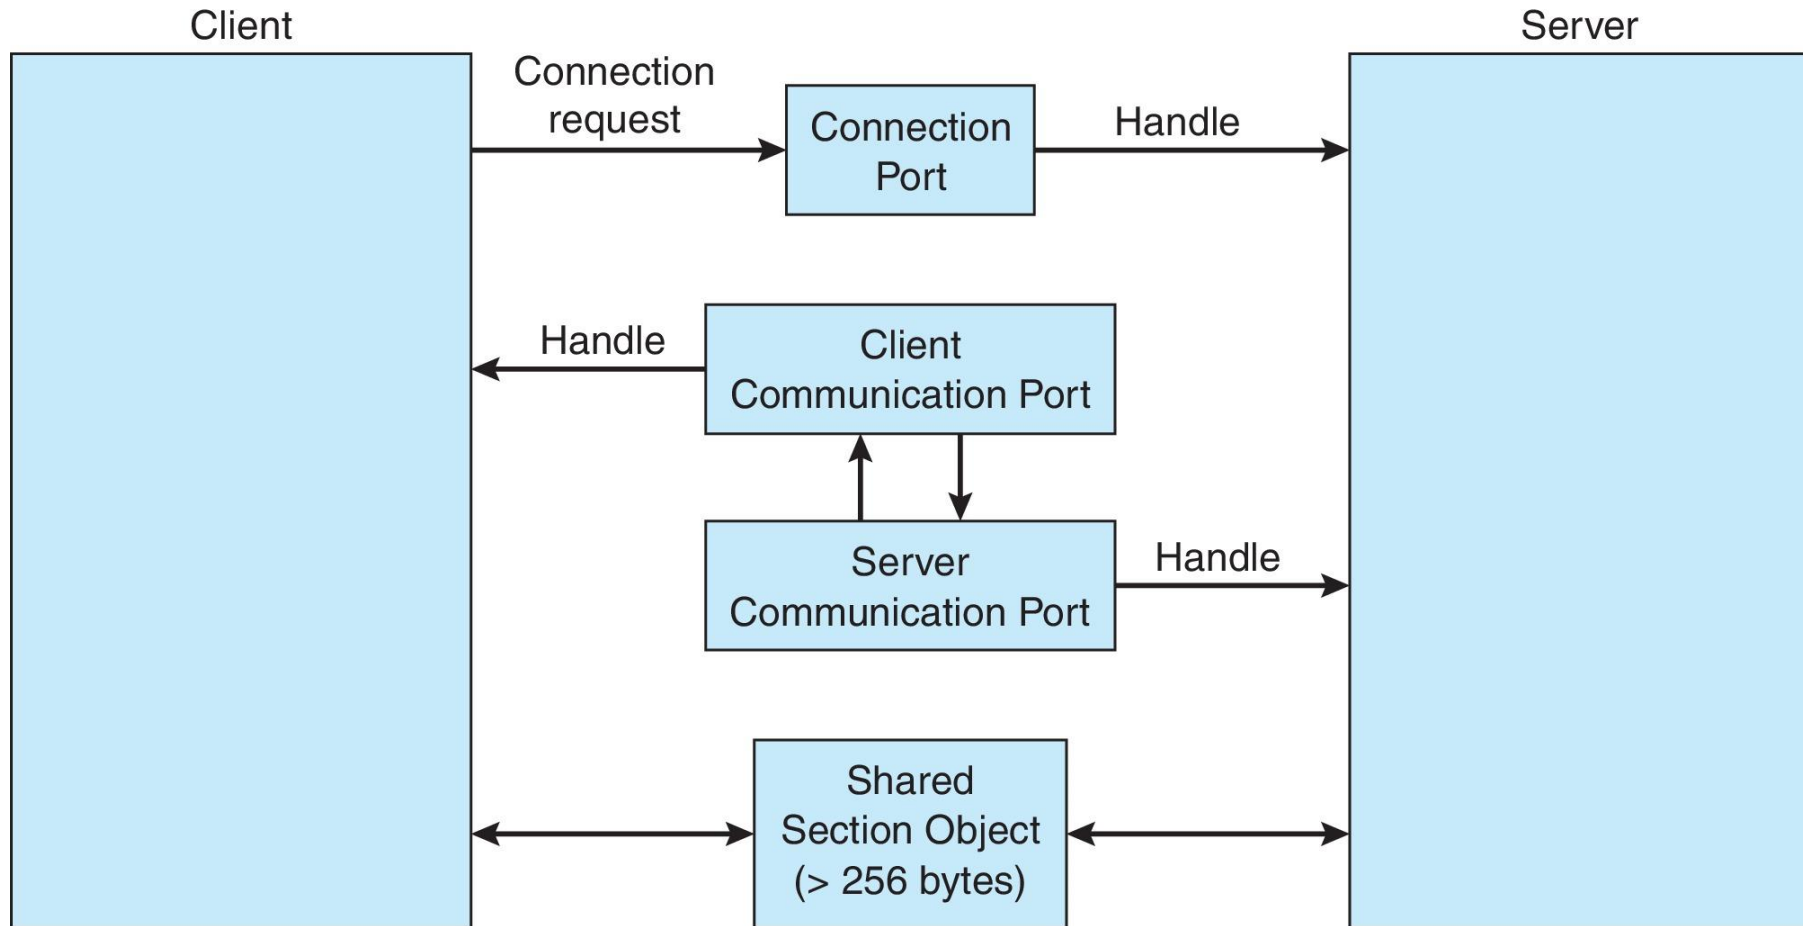
Communication works as follows:

- ▶ The client opens a handle to the subsystem's **connection port** object.
- ▶ The client sends a connection request.
- ▶ The server creates two private **communication ports** and returns the handle to one of them to the client.
- ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.





Local Procedure Calls in Windows





Pipes

Acts as a conduit allowing two processes to communicate

Issues:

Is communication unidirectional or bidirectional?

In the case of two-way communication, is it half or full-duplex?

Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?

Can the pipes be used over a network?

Ordinary pipes – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.

Named pipes – can be accessed without a parent-child relationship.





Ordinary Pipes


Ordinary Pipes allow communication in standard producer-consumer style

Producer writes to one end (the **write-end** of the pipe)

Consumer reads from the other end (the **read-end** of the pipe)

Ordinary pipes are therefore unidirectional

Require parent-child relationship between communicating processes



Windows calls these **anonymous pipes**





Named Pipes

Named Pipes are more powerful than ordinary pipes

Communication is bidirectional

No parent-child relationship is necessary between the communicating processes

Several processes can use the named pipe for communication

Provided on both UNIX and Windows systems





Communications in Client-Server Systems

Sockets

Remote Procedure Calls





Sockets

A **socket** is defined as an endpoint for communication

Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host

The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**

Communication consists between a pair of sockets

All ports below 1024 are ***well known***, used for standard services

Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running







Socket Communication

host X
(146.86.5.20)



web server
(161.25.19.8)





Sockets in Java

Three types of sockets

Connection-oriented
(TCP)

Connectionless
(UDP)

MulticastSocket
class— data can be
sent to multiple
recipients

Consider this “Date”
server in Java:

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





Sockets in Java

The equivalent Date client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





Remote Procedure Calls

Remote procedure call (RPC) abstracts procedure calls between processes on networked systems

Again uses ports for service differentiation

Stubs – client-side proxy for the actual procedure on the server

The client-side stub locates the server and **marshalls** the parameters

The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server

On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**





Remote Procedure Calls (Cont.)

Data representation handled via **External Data Representation (XDL)** format to account for different architectures

Big-endian and **little-endian**

Remote communication has more failure scenarios than local

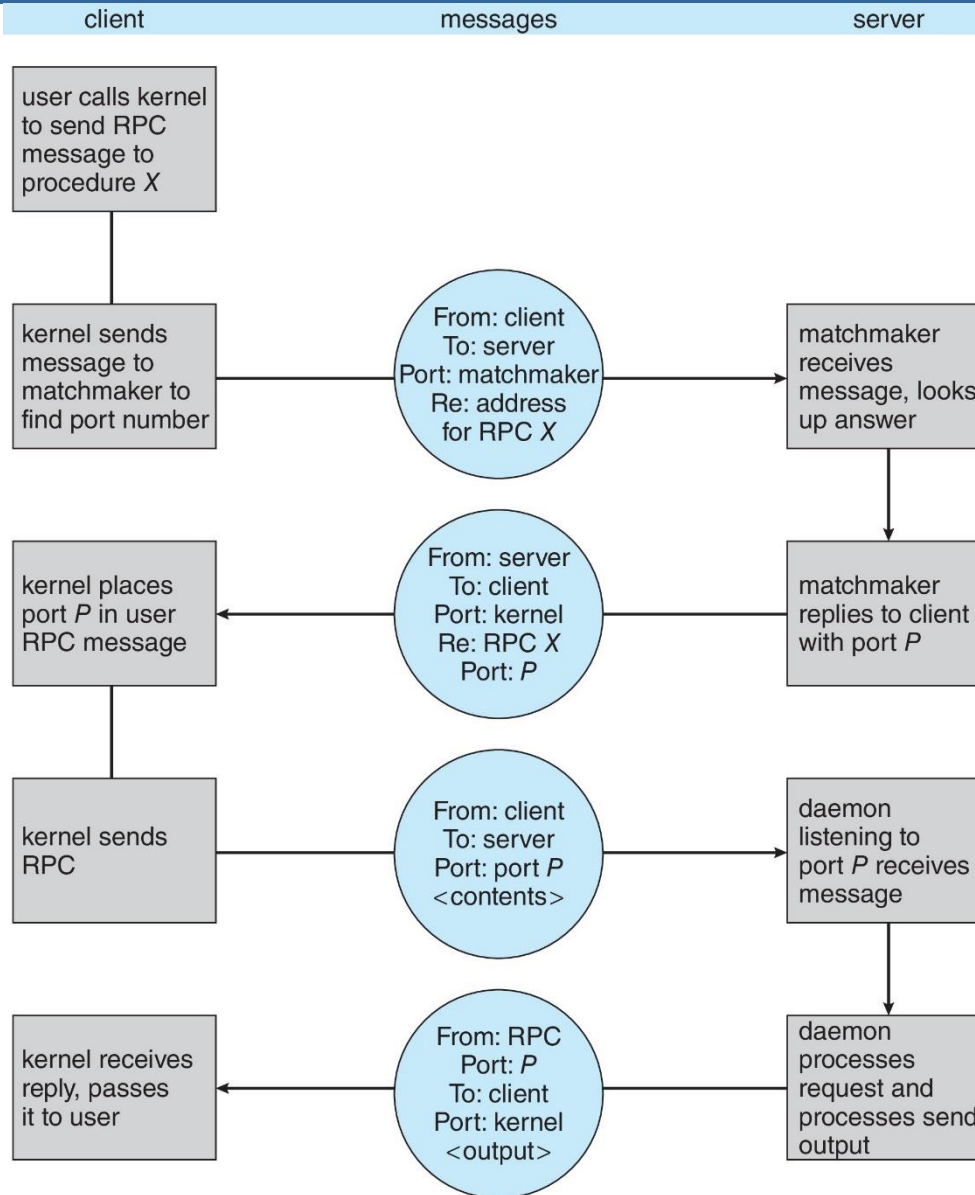
Messages can be delivered ***exactly once*** rather than ***at most once***

OS typically provides a rendezvous (or **matchmaker**) service to connect client and server





Execution of RPC



End of Chapter 3

