


# Chapter 4: Threads & Concurrency

---





# Chapter 4: Threads

---

Overview

Multicore Programming

Multithreading Models

Thread Libraries

Implicit Threading

Threading Issues

Operating System Examples





# Objectives

---

Identify the basic components of a thread, and contrast threads and processes

Describe the benefits and challenges of designing multithreaded applications

Illustrate different approaches to implicit threading including thread pools, fork-join, and Grand Central Dispatch

Describe how the Windows and Linux operating systems represent threads

Design multithreaded applications using the Pthreads, Java, and Windows threading APIs





# Motivation

---

Most modern applications are multithreaded

Threads run within application

Multiple tasks with the application can be implemented by separate threads

- Update display

- Fetch data

- Spell checking

- Answer a network request

Process creation is heavy-weight while thread creation is light-weight

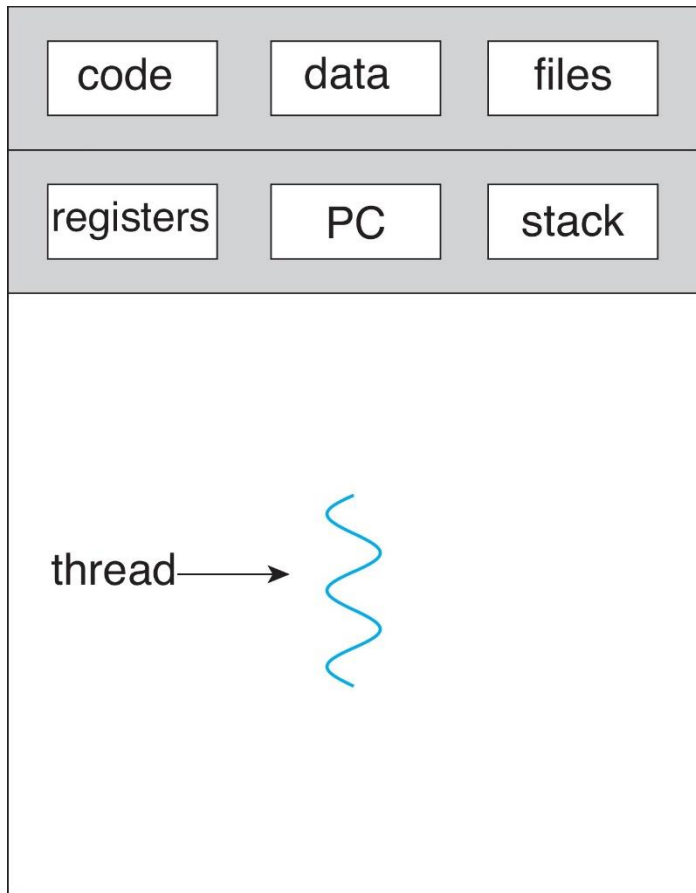
Can simplify code, increase efficiency

Kernels are generally multithreaded






# Single and Multithreaded Processes



single-threaded process

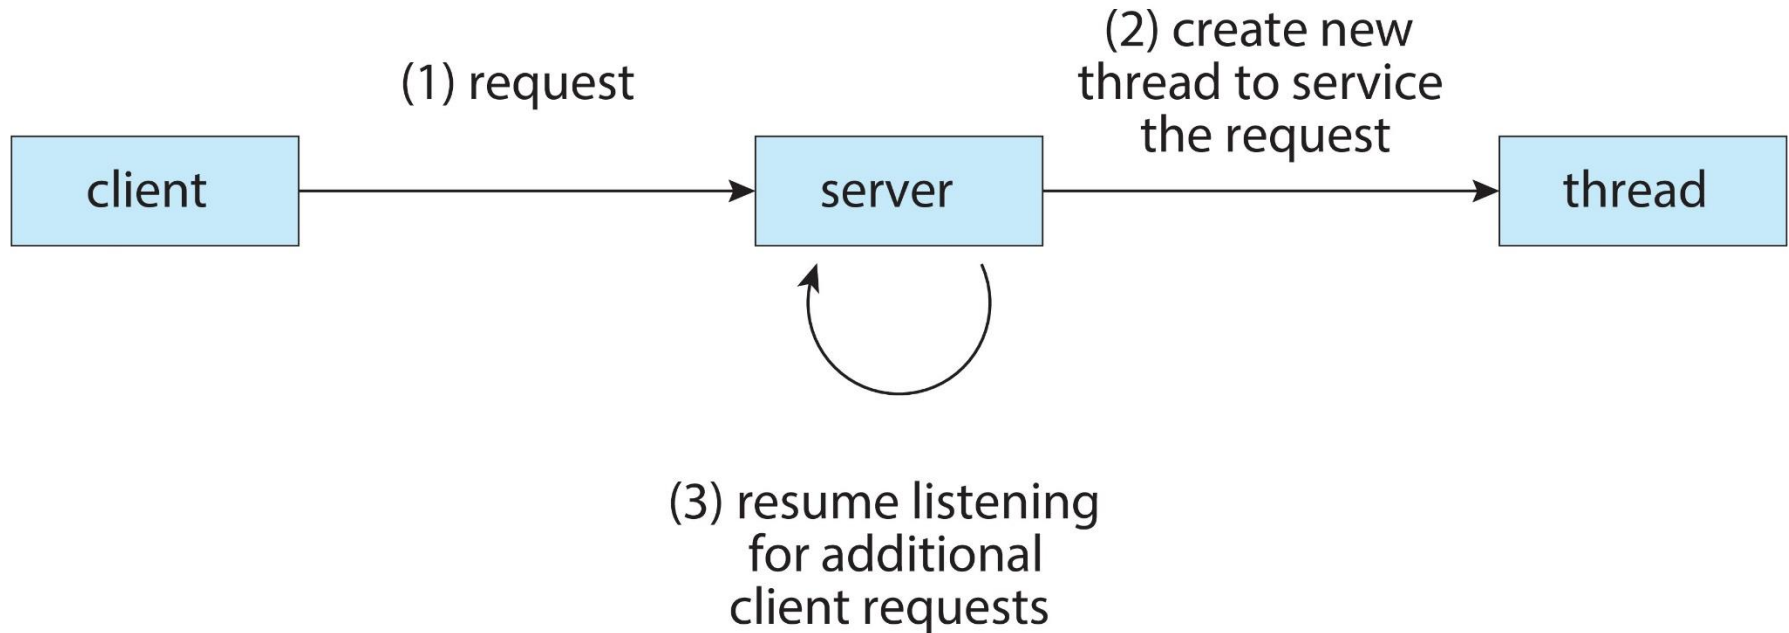


multithreaded process





# Multithreaded Server Architecture





# Benefits

---

**Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces

**Resource Sharing** – threads share resources of process, easier than shared memory or message passing

**Economy** – cheaper than process creation, thread switching lower overhead than context switching

**Scalability** – process can take advantage of multicore architectures





# Multicore Programming

---

**Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:

**Dividing activities**

**Balance**

**Data splitting**

**Data dependency**

**Testing and debugging**

***Parallelism*** implies a system can perform more than one task simultaneously

***Concurrency*** supports more than one task making progress

Single processor / core, scheduler providing concurrency









# Concurrency vs. Parallelism

## Concurrent execution on single-core system:



## Parallelism on a multi-core system:





# Multicore Programming

---

## Types of parallelism



**Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each

**Task parallelism** – distributing threads across cores, each thread performing unique operation





# Data and Task Parallelism





# Amdahl's Law

Identifies performance gains from adding additional cores to an application that has both serial and parallel components

$S$  is serial portion

$N$  processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times

As  $N$  approaches infinity, speedup approaches  $1 / S$


**Serial portion of an application has disproportionate effect on performance gained by adding additional cores**

But does the law take into account contemporary multicore systems?





# Amdahl's Law





# User Threads and Kernel Threads

---

**User threads** - management done by user-level threads library

Three primary thread libraries:

- POSIX **Pthreads**

- Windows threads

- Java threads

**Kernel threads** - Supported by the Kernel

Examples – virtually all general purpose operating systems, including:

- Windows

- Linux

- Mac OS X


- iOS

- Android





# User and Kernel Threads






# Multithreading Models

---

Many-to-One

One-to-One

Many-to-Many







# Many-to-One

Many user-level threads mapped to single kernel thread

One thread blocking causes all to block


Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time

Few systems currently use this model

Examples:

**Solaris Green Threads**

**GNU Portable Threads**





# One-to-One

Each user-level thread maps to kernel thread

Creating a user-level thread creates a kernel thread


More concurrency than many-to-one

Number of threads per process sometimes restricted due to overhead

Examples

Windows

Linux






# Many-to-Many Model

Allows many user level threads to be mapped to many kernel threads

Allows the operating system to create a sufficient number of kernel threads

Windows with the *ThreadFiber* package


Otherwise not very common





# Two-level Model

Similar to M:M, except that it allows a user thread to be **bound** to kernel thread





# Thread Libraries

---

**Thread library** provides programmer with API for creating and managing threads

Two primary ways of implementing

- Library entirely in user space

- Kernel-level library supported by the OS





# Pthreads

---

May be provided either as user-level or kernel-level

A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization

***Specification***, not ***implementation***

API specifies behavior of the thread library,  
implementation is up to development of the library

Common in UNIX operating systems (Linux & Mac OS X)





# Pthreads Example

---

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```





# Pthreads Example (cont)

---

```
/* The thread will execute in this function */  
void *runner(void *param)  
{  
    int i, upper = atoi(param);  
    sum = 0;  
  
    for (i = 1; i <= upper; i++)  
        sum += i;  
  
    pthread_exit(0);  
}
```







# Pthreads Code for Joining 10 Threads

---

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```





# Windows Multithreaded C Program

---

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
```





# Windows Multithreaded C Program (Cont.)

```
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```





# Java Threads

---

Java threads are managed by the JVM

Typically implemented using the threads model provided by underlying OS

Java threads may be created by:

- Extending Thread class

- Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```

Standard practice is to implement Runnable interface





# Java Threads

## Implementing Runnable interface:

```
class Task implements Runnable
{
    public void run() {
        System.out.println("I am a thread.");
    }
}
```

## Creating a thread:

```
Thread worker = new Thread(new Task());
worker.start();
```

## Waiting on a thread:

```
try {
    worker.join();
}
catch (InterruptedException ie) { }
```





# Java Executor Framework

---

Rather than explicitly creating threads, Java also allows thread creation around the Executor interface:

```
public interface Executor
{
    void execute(Runnable command);
}
```

The Executor is used as follows:

```
Executor service = new Executor();
service.execute(new Task());
```





# Java Executor Framework

---

```
import java.util.concurrent.*;

class Summation implements Callable<Integer>
{
    private int upper;
    public Summation(int upper) {
        this.upper = upper;
    }

    /* The thread will execute in this method */
    public Integer call() {
        int sum = 0;
        for (int i = 1; i <= upper; i++)
            sum += i;

        return new Integer(sum);
    }
}
```





# Java Executor Framework (cont)

---

```
public class Driver
{
    public static void main(String[] args) {
        int upper = Integer.parseInt(args[0]);

        ExecutorService pool = Executors.newSingleThreadExecutor();
        Future<Integer> result = pool.submit(new Summation(upper));

        try {
            System.out.println("sum = " + result.get());
        } catch (InterruptedException | ExecutionException ie) { }
    }
}
```







# Implicit Threading

---

Growing in popularity as numbers of threads increase,  
program correctness more difficult with explicit threads

Creation and management of threads done by compilers and  
run-time libraries rather than programmers

Five methods explored

- Thread Pools

- Fork-Join

- OpenMP

- Grand Central Dispatch

- Intel Threading Building Blocks





# Thread Pools

Create a number of threads in a pool where they await work

Advantages:

Usually slightly faster to service a request with an existing thread than create a new thread

Allows the number of threads in the application(s) to be bound to the size of the pool

Separating task to be performed from mechanics of creating task allows different strategies for running task

- ▶ i.e. Tasks could be scheduled to run periodically

Windows API supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
    * this function runs as a separate thread.  
    */  
}
```





# Java Thread Pools

---

Three factory methods for creating thread pools in Executors class:

- `static ExecutorService newSingleThreadExecutor()`
- `static ExecutorService newFixedThreadPool(int size)`
- `static ExecutorService newCachedThreadPool()`





# Java Thread Pools (cont)

---

```
import java.util.concurrent.*;

public class ThreadPoolExample
{
    public static void main(String[] args) {
        int numTasks = Integer.parseInt(args[0].trim());

        /* Create the thread pool */
        ExecutorService pool = Executors.newCachedThreadPool();

        /* Run each task using a thread in the pool */
        for (int i = 0; i < numTasks; i++)
            pool.execute(new Task());


        /* Shut down the pool once all threads have completed */
        pool.shutdown();
    }
}
```





# Fork-Join Parallelism

Multiple threads (tasks) are **forked**, and then **joined**.





# Fork-Join Parallelism

---

General algorithm for fork-join strategy:

```
Task(problem)
  if problem is small enough
    solve the problem directly
  else
    subtask1 = fork(new Task(subset of problem))
    subtask2 = fork(new Task(subset of problem))


    result1 = join(subtask1)
    result2 = join(subtask2)

    return combined results
```





# Fork-Join Parallelism





# Fork-Join Parallelism in Java

---

```
ForkJoinPool pool = new ForkJoinPool();  
// array contains the integers to be summed  
int[] array = new int[SIZE];  
  
SumTask task = new SumTask(0, SIZE - 1, array);  
int sum = pool.invoke(task);
```







# Fork-Join Parallelism in Java

```
import java.util.concurrent.*;

public class SumTask extends RecursiveTask<Integer>
{
    static final int THRESHOLD = 1000;

    private int begin;
    private int end;
    private int[] array;

    public SumTask(int begin, int end, int[] array) {
        this.begin = begin;
        this.end = end;
        this.array = array;
    }

    protected Integer compute() {
        if (end - begin < THRESHOLD) {
            int sum = 0;
            for (int i = begin; i <= end; i++)
                sum += array[i];

            return sum;
        }
        else {
            int mid = (begin + end) / 2;

            SumTask leftTask = new SumTask(begin, mid, array);
            SumTask rightTask = new SumTask(mid + 1, end, array);

            leftTask.fork();
            rightTask.fork();

            return rightTask.join() + leftTask.join();
        }
    }
}
```





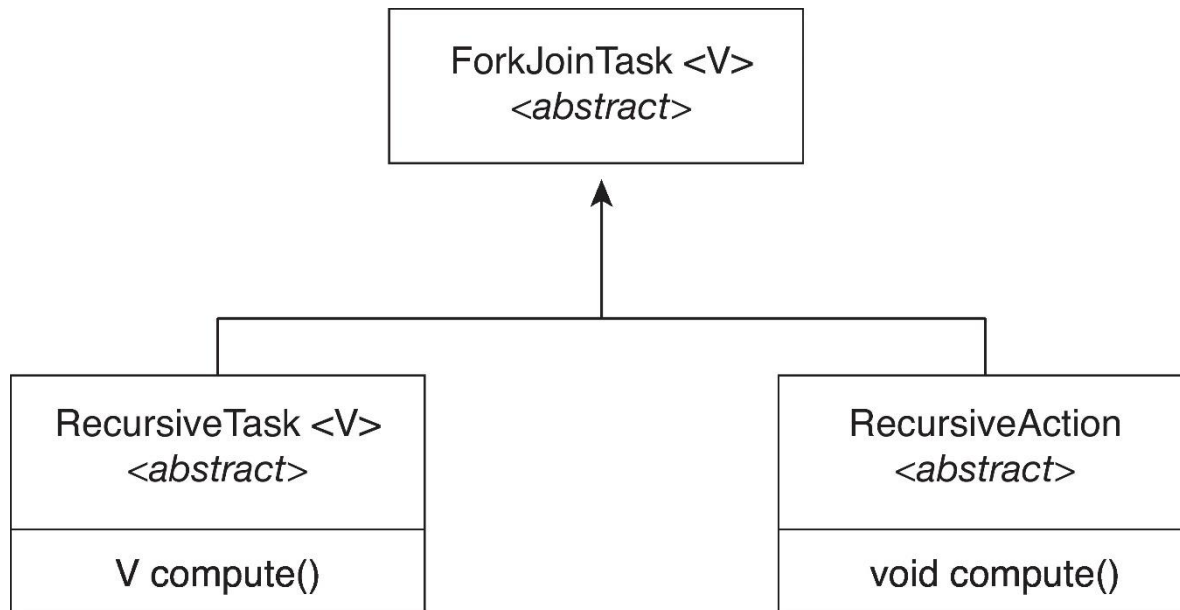
# Fork-Join Parallelism in Java

The **ForkJoinTask** is an abstract base class

**RecursiveTask** and **RecursiveAction** classes extend **ForkJoinTask**

**RecursiveTask** returns a result (via the return value from the `compute()` method)

**RecursiveAction** does not return a result





# OpenMP

Set of compiler directives  
and an API for C, C++,  
FORTRAN

Provides support for parallel  
programming in shared-  
memory environments

Identifies **parallel regions** –  
blocks of code that can run  
in parallel

**#pragma omp parallel**

Create as many threads as  
there are cores

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```





# OpenMP

---

Run the for loop in parallel

```
#pragma omp parallel for  
for (i = 0; i < N; i++) {  
    c[i] = a[i] + b[i];  
}
```





# Grand Central Dispatch

---

Apple technology for macOS and iOS operating systems

Extensions to C, C++ and Objective-C languages, API, and run-time library

Allows identification of parallel sections

Manages most of the details of threading

Block is in “ $\wedge\{ \}$ ” :

```
 $\wedge\{ \text{ printf("I am a block"); } \}$ 
```

Blocks placed in dispatch queue

Assigned to available thread in thread pool when removed from queue





# Grand Central Dispatch

---

Two types of dispatch queues:

**serial** – blocks removed in FIFO order, queue is per process, called **main queue**

- ▶ Programmers can create additional serial queues within program

**concurrent** – removed in FIFO order but several may be removed at a time

- ▶ Four system wide queues divided by quality of service:

- `QOS_CLASS_USER_INTERACTIVE`
- `QOS_CLASS_USER_INITIATED`
- `QOS_CLASS_USER_UTILITY`
- `QOS_CLASS_USER_BACKGROUND`





# Grand Central Dispatch

---

For the Swift language a task is defined as a closure – similar to a block, minus the caret

Closures are submitted to the queue using the `dispatch_async()` function:

```
let queue = dispatch_get_global_queue  
            (QOS_CLASS_USER_INITIATED, 0)
```

```
dispatch_async(queue, { print("I am a closure.") })
```





# Intel Threading Building Blocks (TBB)

---

Template library for designing parallel C++ programs

A serial version of a simple for loop

```
for (int i = 0; i < n; i++) {  
    apply(v[i]);  
}
```

The same for loop written using TBB with **parallel\_for** statement:

```
parallel_for (size_t(0), n, [=](size_t i) {apply(v[i]);});
```







# Threading Issues

---

Semantics of **fork()** and **exec()** system calls

Signal handling

Synchronous and asynchronous

Thread cancellation of target thread

Asynchronous or deferred

Thread-local storage

Scheduler Activations





# Semantics of `fork()` and `exec()`

---

Does `fork()` duplicate only the calling thread or all threads?

Some UNIXes have two versions of `fork`

`exec()` usually works as normal – replace the running process including all threads





# Signal Handling

- n **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- n A **signal handler** is used to process signals
  1. Signal is generated by particular event
  2. Signal is delivered to a process
  3. Signal is handled by one of two signal handlers:
    1. default
    2. user-defined
- n Every signal has **default handler** that kernel runs when handling signal
  - | **User-defined signal handler** can override default
  - | For single-threaded, signal delivered to process





# Signal Handling (Cont.)

---

- n Where should a signal be delivered for multi-threaded?
  - | Deliver the signal to the thread to which the signal applies
  - | Deliver the signal to every thread in the process
  - | Deliver the signal to certain threads in the process
  - | Assign a specific thread to receive all signals for the process





# Thread Cancellation

---

Terminating a thread before it has finished

Thread to be canceled is **target thread**

Two general approaches:

**Asynchronous cancellation** terminates the target thread immediately

**Deferred cancellation** allows the target thread to periodically check if it should be cancelled

Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```





# Thread Cancellation (Cont.)

Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

| Mode         | State    | Type         |
|--------------|----------|--------------|
| Off          | Disabled | —            |
| Deferred     | Enabled  | Deferred     |
| Asynchronous | Enabled  | Asynchronous |


If thread has cancellation disabled, cancellation remains pending until thread enables it

Default type is deferred

Cancellation only occurs when thread reaches **cancellation point**

- ▶ I.e. `pthread_testcancel()`
- ▶ Then **cleanup handler** is invoked

On Linux systems, thread cancellation is handled through signals





# Thread Cancellation in Java

---

Deferred cancellation uses the `interrupt()` method, which sets the interrupted status of a thread.

```
Thread worker;  
  
. . .  
  
/* set the interruption status of the thread */  
worker.interrupt()
```

A thread can then check to see if it has been interrupted:

```
while (!Thread.currentThread().isInterrupted()) {  
    . . .  
}
```





# Thread-Local Storage

---

**Thread-local storage (TLS)** allows each thread to have its own copy of data

Useful when you do not have control over the thread creation process (i.e., when using a thread pool)

Different from local variables

- Local variables visible only during single function invocation

- TLS visible across function invocations

Similar to **static** data

- TLS is unique to each thread







# Scheduler Activations

Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application

Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**


Appears to be a virtual processor on which process can schedule user thread to run

Each LWP attached to kernel thread

How many LWPs to create?

Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library

This communication allows an application to maintain the correct number kernel threads





# Operating System Examples

---

Windows Threads

Linux Threads





# Windows Threads

---

Windows API – primary API for Windows applications

Implements the one-to-one mapping, kernel-level

Each thread contains

- A thread id

- Register set representing state of processor

- Separate user and kernel stacks for when thread runs in user mode or kernel mode

- Private data storage area used by run-time libraries and dynamic link libraries (DLLs)

The register set, stacks, and private storage area are known as the **context** of the thread





# Windows Threads (Cont.)

---

The primary data structures of a thread include:

ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space

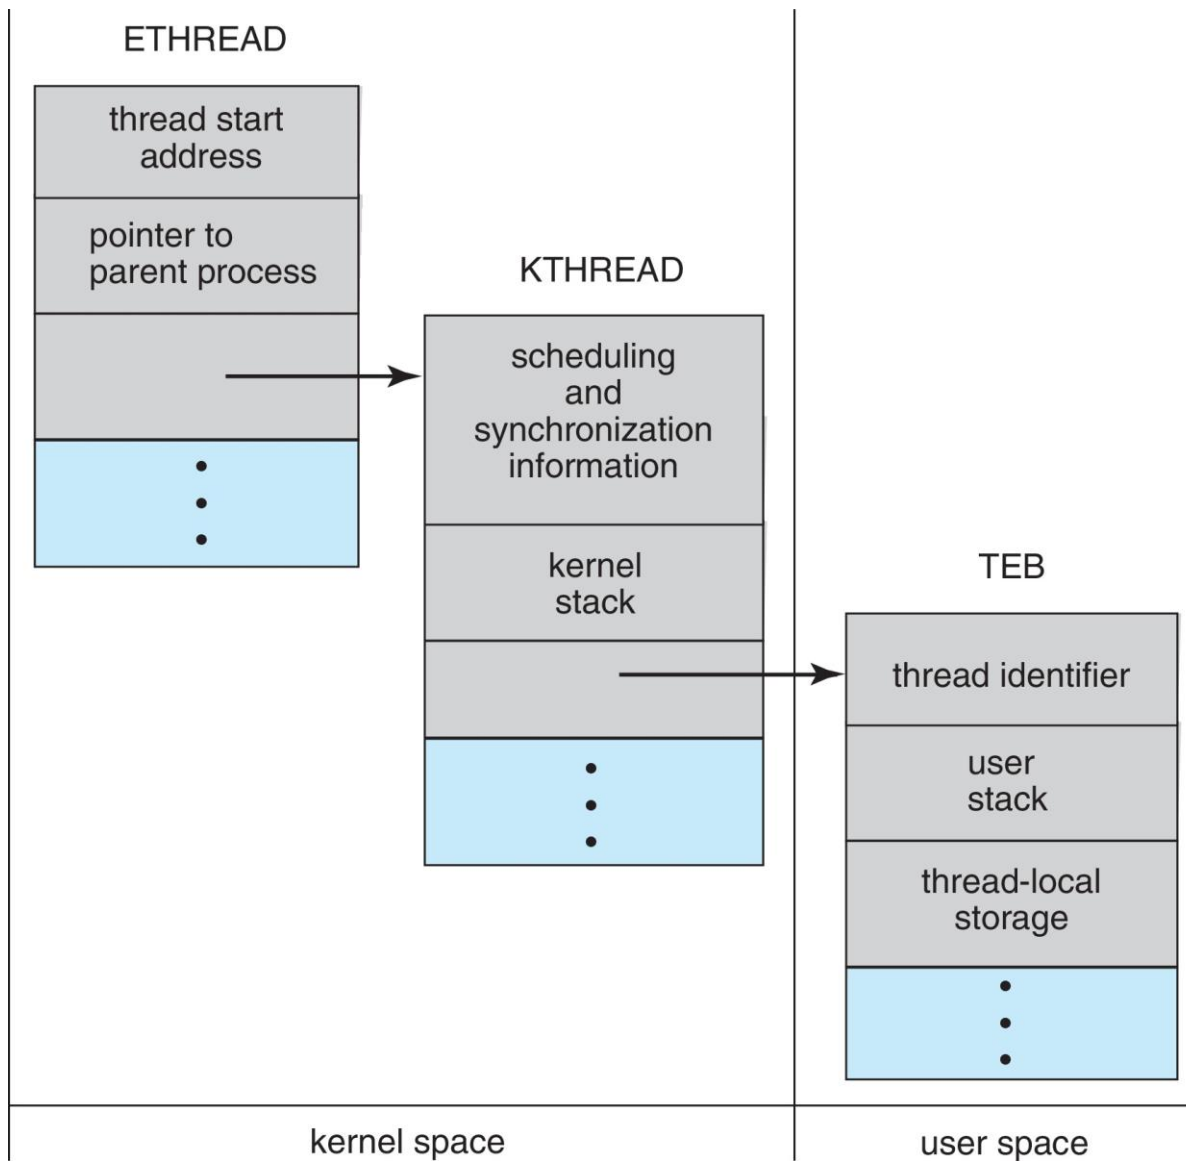
KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space

TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space





# Windows Threads Data Structures





# Linux Threads

Linux refers to them as **tasks** rather than **threads**

Thread creation is done through `clone()` system call

`clone()` allows a child task to share the address space of the parent task (process)

Flags control behavior

| flag          | meaning                            |
|---------------|------------------------------------|
| CLONE_FS      | File-system information is shared. |
| CLONE_VM      | The same memory space is shared.   |
| CLONE_SIGHAND | Signal handlers are shared.        |
| CLONE_FILES   | The set of open files is shared.   |

`struct task_struct` points to process data structures (shared or unique)



# End of Chapter 4

---

